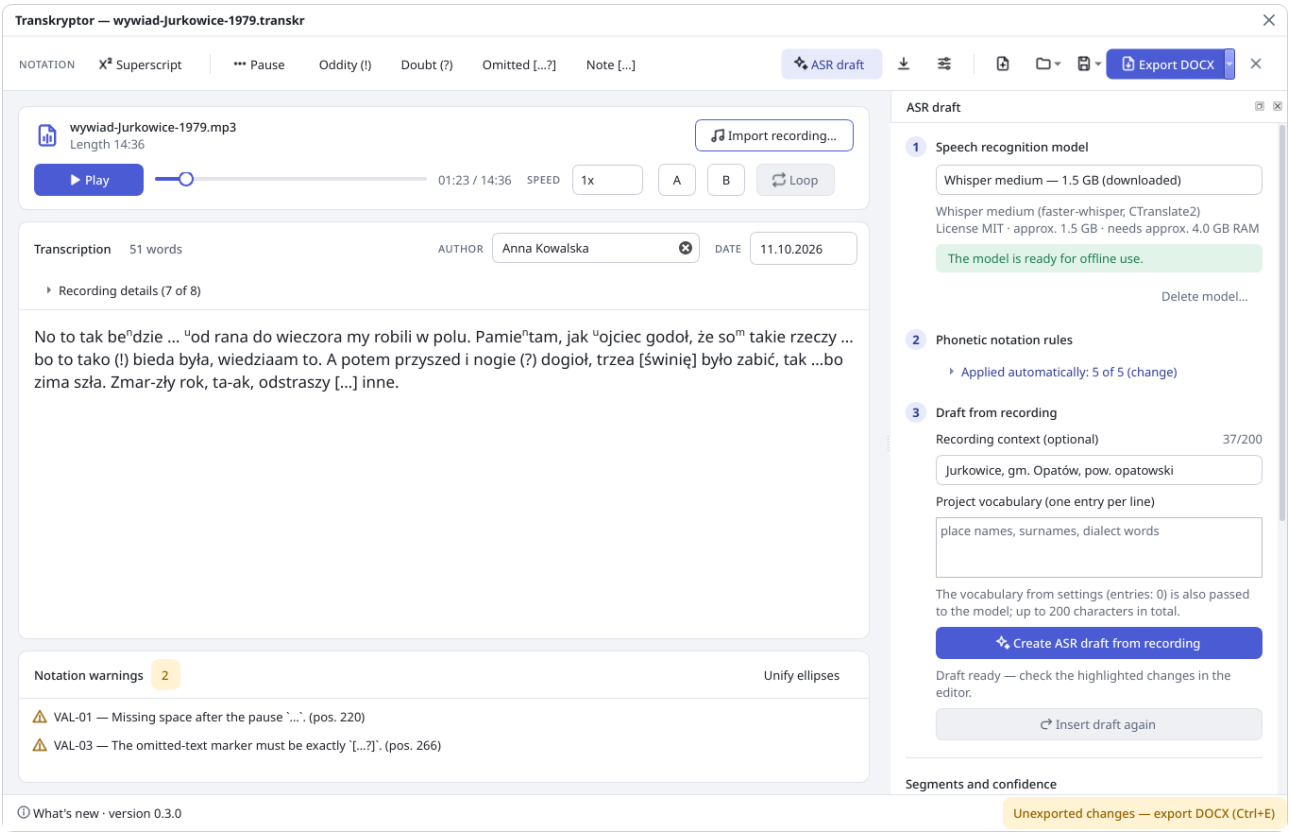
TRANSKRYPTOR

User guide

A local desktop application for manual phonetic transcription of Polish speech, assisted by an optional automatic speech recognition draft.



Transkryptor at work

Application version 0.3.0 · Windows 10/11 (64-bit) and Linux x86_64 · October 2026

Contents

1. What Transkryptor is 3

2. Requirements 3

3. Installation 4

4. The application window 5

5. Buttons and keyboard shortcuts 7

6. Working with the recording and the text 9

7. Recording details 13

8. ASR draft (automatic speech recognition) 14

9. Project and autosave 18

10. Export to DOCX 19

11. DOCX import 20

12. Closing and unsaved changes 21

13. Privacy 21

14. Troubleshooting 22

15. Uninstalling 23

Appendix: notation cheat sheet 24

1. What Transkryptor is

Transkryptor helps you write down a recorded utterance the way it sounds, not the way it is spelled. You listen to an MP3 recording, type the transcription in the editor, mark pronunciation phenomena (nasality, an inserted *u*, pauses, uncertain passages), and export the finished text to a Word file (DOCX).

The application runs entirely on your computer. There is no account, no sign-in, and recordings are never sent to the Internet.

Transkryptor is built for Polish dialect speech, so the notation and the examples in this guide are Polish. The user interface is available in English and Polish — see **Settings → Appearance → Language** (chapter 6).

What it can do

- play an MP3 recording with adjustable speed (0.5×–2×) and seeking,
- insert notation markers with a single shortcut: pause, (!), (?), [...?], [...],
- mark letters as superscript (e.g. be^{n}dzie, ^{u}od),
- warn about notation errors as you type (e.g. an unclosed bracket),
- export the transcription together with the author, date and recording details to DOCX, keeping superscripts — also in an anonymized version,
- save your work in a project file and come back to it later; autosave protects it against crashes,
- load a previously exported (or someone else's) DOCX document as a new project,
- optionally: create an automatic text draft (ASR), insert it straight into the editor and apply selected phonetic notation rules — with every change highlighted for your review,
- follow the light or dark theme of your system.

> **Important.** You save your work in a project file (**Save project**, `Ctrl+S`) — with the text, recording details, ASR draft and the position in the recording. The DOCX file is the result for other people. Autosave is only a safety net in case of a crash; it does not replace saving the project (chapter 9).

2. Requirements

	Minimum	Recommended
System	Windows 10/11 64-bit or Linux x86_64 (glibc 2.35+, e.g. Ubuntu 22.04 or later)	
RAM	4 GB	8 GB
Disk space	approx. 0.5 GB (application only); 3 GB with the ASR model	
Processor	4 cores	8 cores or more
Internet	only once, to download the ASR model (1.5 GB)	

You do not need Python or any other tools — the installer contains everything. On Linux, playback needs working system audio (PipeWire or PulseAudio; typical distributions have it by default).

Manual work (typing, listening, export) runs on much weaker hardware. The requirements mainly concern the ASR draft. On the test machine an hour of recording was processed in about half an hour; a slower processor will take longer.

3. Installation

Download the file for your system and the **SHA256SUMS.txt** file.

System	File
Windows	Transkryptor-0.3.0-windows-x64-setup.exe
Linux	Transkryptor-0.3.0-x86_64.AppImage

Windows

1. Double-click **Transkryptor-0.3.0-windows-x64-setup.exe**.
2. If a **Windows protected your PC** window (SmartScreen) appears, click **More info** → **Run anyway**. The application is not signed with a commercial certificate, hence the warning.
3. By default the installer installs the application for your account only and does not need administrator rights. In the first window you can choose to install it for all users instead.
4. Accept the license (GPL-3.0), check the desktop shortcut option if you want one, and click **Install**.
5. At the end you can start the application right away. Later you will find it in **Start menu → Transkryptor**.

Linux (AppImage)

An AppImage is a single file that needs neither installation nor administrator rights.

1. Save the file, e.g. in the `~/Applications` folder.
2. Make it executable — in the file manager: **Properties → Permissions → Allow executing**, or in a terminal:

```
chmod +x Transkryptor-0.3.0-x86_64.AppImage
./Transkryptor-0.3.0-x86_64.AppImage
```

- If the application does not start and a message about FUSE appears, install the `libfuse2` package (Ubuntu: `sudo apt install libfuse2`).
- You can add a system menu entry with a tool such as Gear Lever or AppImageLauncher.

Verifying the downloaded file (optional)

The checksum confirms that the file has not been damaged or tampered with.

```
# Linux
sha256sum --check SHA256SUMS.txt

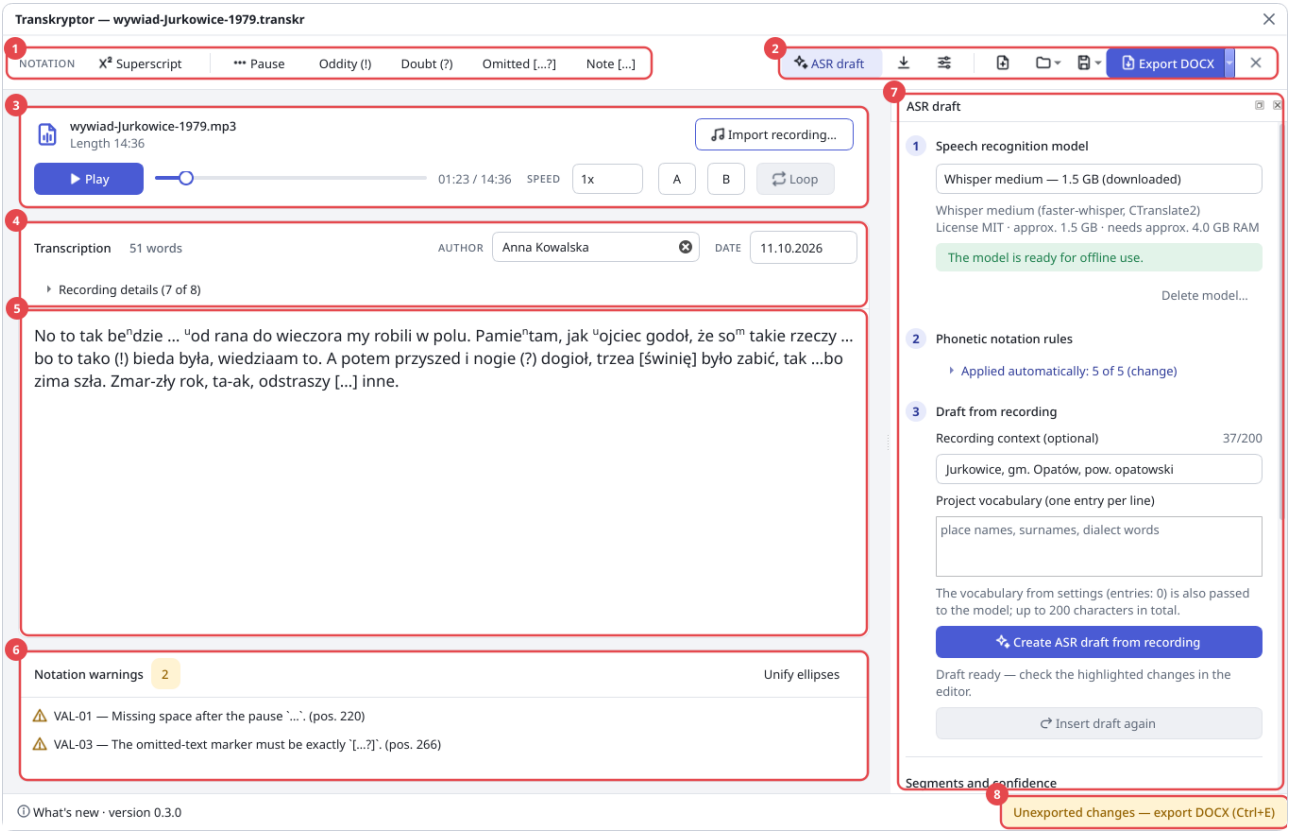
# Windows (PowerShell) — compare the result with SHA256SUMS.txt
```

```
Get-FileHash .\Transkryptor-0.3.0-windows-x64-setup.exe -Algorithm SHA256
```

Diagnostics without opening the window: run the application with the `--self-test` switch. Every item should have the `[ok]` status; the line “ASR model: not downloaded” is not an error.

4. The application window

After starting the application you will see one window: the toolbar at the top, the working area on the left (player, transcription, warnings) and the **ASR draft** panel on the right. Drag the borders between the areas with the mouse to change their proportions.



The main window at work — the numbers match the table below

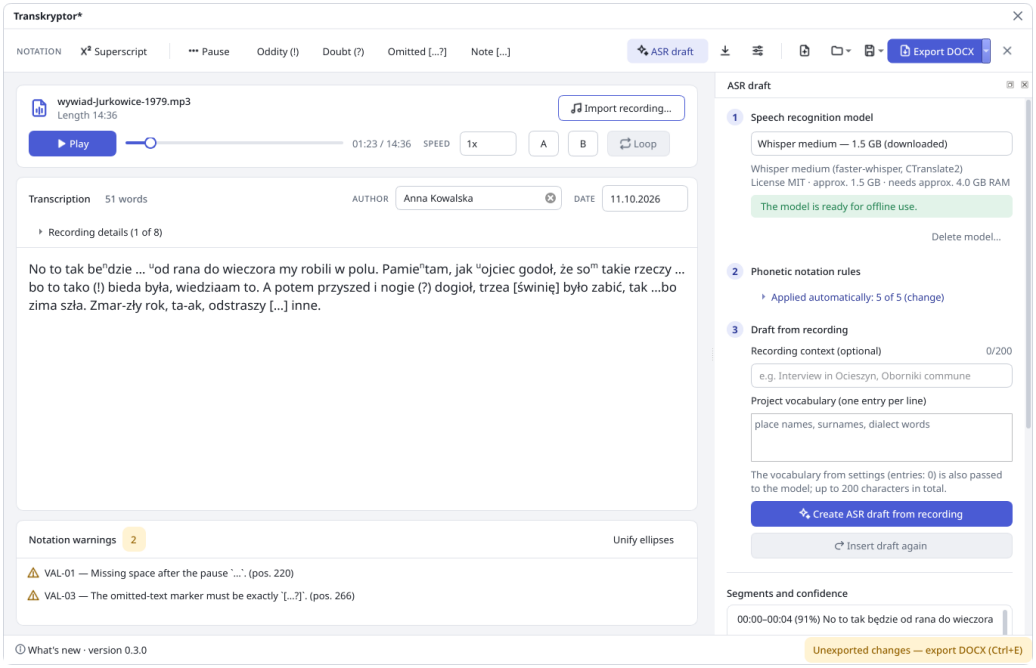
No.	Element	What it is for
1	Toolbar — Notation	superscript and inserting markers (chapter 5)
2	Toolbar — actions	ASR panel, updates, settings, new document, opening and saving the project, export to DOCX (highlighted button), closing the application
3	Player	recording import, name and length, playback, seeking, speed
4	Transcription header	word counter, the Author and Date fields and the collapsible recording details (chapter 7) — saved in the project and in the DOCX file
5	Editor	the place for the transcription text

No.	Element	What it is for
6	Notation warnings	list of detected deviations from the notation, with a counter
7	ASR draft panel	model download, rules, automatic draft, change review (chapter 8)
8	Status bar	on the left “What's new” (the change history), on the right — whether all changes have been exported

> **Tip.** The window title shows the name of the open project. An asterisk (e.g. *Transkryptor — interview.transkr**) means unsaved project changes or, when the work has no project file yet, unexported changes. A yellow message in the status bar reminds you to export to DOCX.

Light and dark theme

By default Transkryptor follows the system theme: when dark mode is selected in Windows or your Linux desktop, the application window is dark too. You can also choose the light or dark theme permanently in **Settings → Appearance**. The look does not affect the exported DOCX file.



The same window in the dark theme

The ASR draft panel

You can detach the panel into a separate window or close it with the buttons in its corner. Bring a closed panel back with the **ASR draft** button in the toolbar.

5. Buttons and keyboard shortcuts

Notation



Toolbar — the Notation group

Button	Shortcut	What it does	Example
Superscript	`Ctrl+Shift+↑`	Turns the selected letters into superscript (or back). Nothing happens without a selection.	be^{n}dzie, ^{u}od
Pause	`Ctrl+.`	Inserts a pause: an ellipsis with a space on both sides, in the form chosen in the settings (… by default).	takich rzeczy … bo
Oddity (!)	`Ctrl+1`	Unusual pronunciation — the spelling is not a mistake.	zamudze (!) pani
Doubt (?)	`Ctrl+2`	Uncertain hearing.	nogie (?) dogioł
Omitted […?]	`Ctrl+3`	Inserts […?] — a fragment left out because it was unintelligible.	odstraszy […?] inne
Note […]	`Ctrl+[`	Inserts […] and puts the cursor inside — type your comment.	trzea [świnię]

Actions



Toolbar — the actions on the right

Button	Shortcut	What it does
ASR draft	—	Shows or hides the ASR draft panel.
Down arrow (updates)	—	Checks whether a newer version of the application is available.
Settings (sliders)	`Ctrl+,`	Opens the settings for appearance and language, player, notation, project and recording details (chapter 6).
New document (sheet with a plus)	`Ctrl+N`	Clears the editor and the metadata and closes the recording. If you have unsaved work, it asks first.
Open project (folder)	`Ctrl+O`	Opens a .transkr project file. The arrow next to it opens a menu: recent projects, Import DOCX… (chapter 11), Clear recent list.
Save project (floppy disk)	`Ctrl+S` `Ctrl+Shift+S`	Saves the project; the first time it asks for a file name. The arrow next to it: Save project as… (chapter 9).

Button	Shortcut	What it does
Export DOCX…	`Ctrl+E`	Saves the transcription to a Word file. The arrow next to it: Export anonymized DOCX… (chapter 10).
× (Close)	`Ctrl+W`	Closes the application (with a warning about unsaved work).

Player

The player shortcuts work while the cursor is in the editor — you can control the recording without taking your hands off the keyboard. They do not insert characters into the text.

Shortcut	Action
`Ctrl+Space` or `F4`	play / pause (after a pause the recording rewinds by 1.5 s — auto-rewind)
`Alt+←` / `Alt+→`	skip the recording back / forward by 3 s (length in the settings)
`Ctrl+Shift+,` / `Ctrl+Shift+.`	slower / faster (next speed from the list)
`Ctrl+Shift+A` / `Ctrl+Shift+B`	loop start (A) / end (B) at the current position
`Ctrl+Shift+L`	turn the A–B loop on / turn it off and clear the markers

Ctrl+Space on Linux. Some input methods (e.g. IBus, fcitx) capture `Ctrl+Space`. Use `F4` instead.

Find and replace

Shortcut	Action
`Ctrl+F`	opens the search bar below the editor
`Ctrl+H`	opens the bar with the “Replace with” field
`F3` / `Shift+F3`	next / previous match
`Esc`	closes the bar and returns to the editor

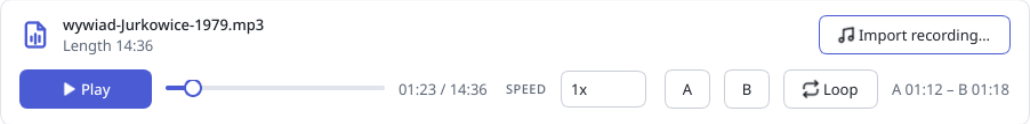
Standard editor shortcuts

Shortcut	Action
`Ctrl+Z` / `Ctrl+Shift+Z`	undo / redo — works for every change, including an inserted ASR draft and applied rules
`Ctrl+C`, `Ctrl+X`, `Ctrl+V`	copy, cut, paste
`Ctrl+A`	select all text

Tip. Hover the mouse over any button to see its description and shortcut. The marker shortcuts work while the cursor is in the editor — you do not need to take your hands off the keyboard while listening.

6. Working with the recording and the text

Player



The player card with a recording loaded

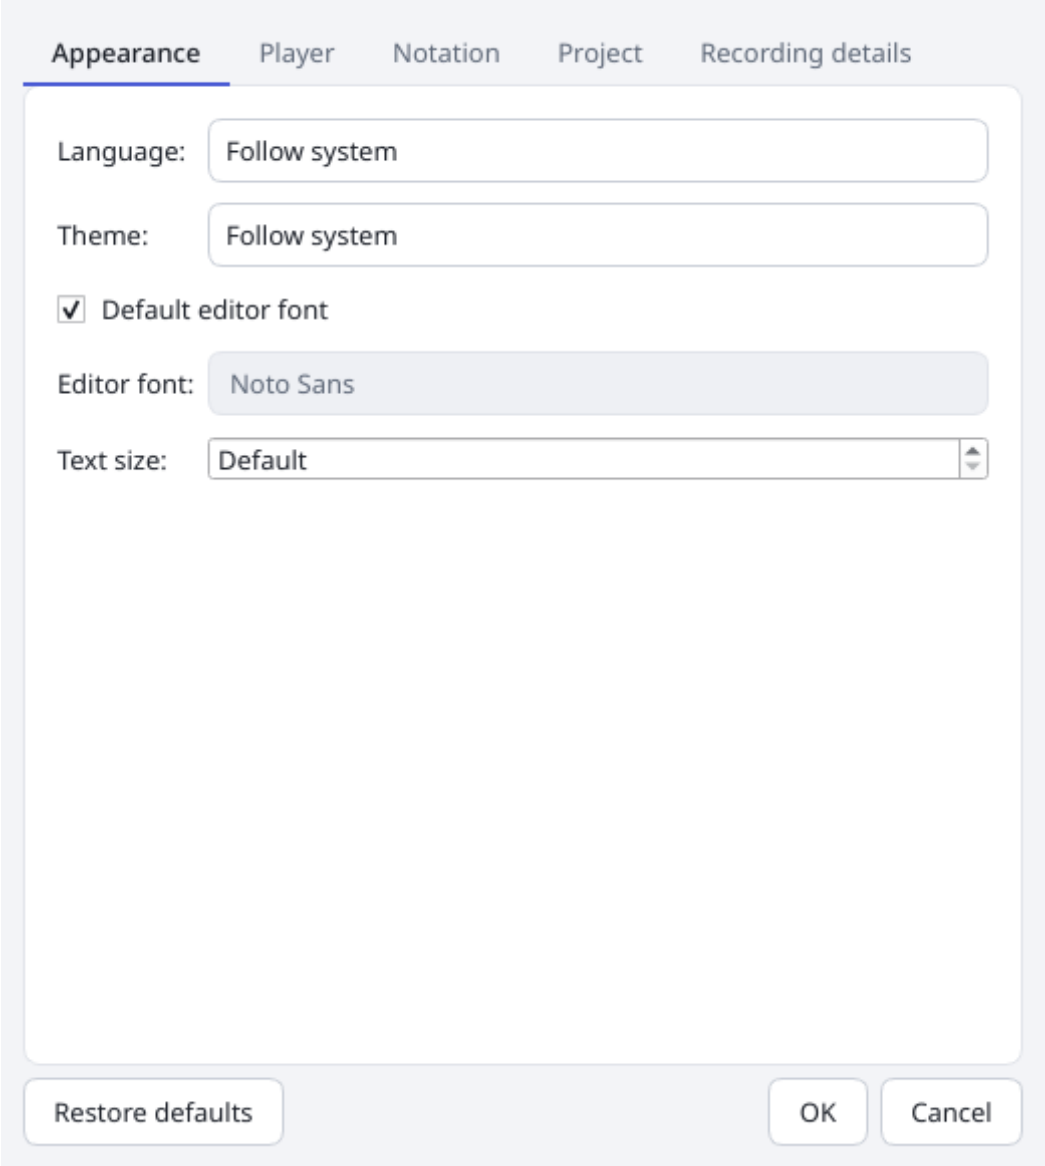
Element	Action
Import recording…	Opens a file dialog (MP3, WAV, FLAC, M4A, AAC, OGG, OPUS, WMA). The file is not copied — the application reads it from where it is.
Recording name	Name of the loaded file and length of the recording. Hover over it to see the full path.
Play / Pause	Starts or pauses playback (the label and icon change with the state).
Slider	Shows the position. Drag it to seek.
Time	Current position / length of the recording (mm:ss).
Speed	0.5x, 0.75x, 1x, 1.25x, 1.5x, 2x — a slower speed makes difficult passages easier to hear.
A, B	Mark the start and end of a fragment to loop at the current position. The range is shown next to them, e.g. A 01:12 – B 01:18.
Loop	Plays the A–B fragment over and over (without B — to the end of the recording). Clicking again turns the loop off and clears the markers.

Loop and auto-rewind

The easiest way to listen to an unclear passage several times is a loop: at its start press `Ctrl+Shift+A`, at its end `Ctrl+Shift+B`, and then `Ctrl+Shift+L`. When playback reaches B it jumps back to A until you turn the loop off.

When you pause playback and resume it, the recording rewinds by 1.5 s so that you hear the context of the last words. Auto-rewind does not apply to the very first start or after seeking during a pause. You can change the rewind length in the settings or turn it off completely.

Settings



The Settings window

Setting	Default	Meaning
Language	follow system	User interface language: English or Polish (Polski). The change takes effect after you restart the application. The DOCX export uses the interface language.
Theme	follow system	Light, dark or the same as the system.
Skip forward and back	3 s	Length of the `Alt+←` / `Alt+→` skip.
Auto-rewind	on, 1.5 s	How far the recording rewinds when resuming after a pause.
Start before ASR segment	0.5 s	A clicked draft segment starts slightly earlier so that the first syllable is not cut off.
Editor font and text size	default	Look of the text in the editor. Does not affect the DOCX export.

Setting	Default	Meaning
Ellipsis style	… (single character)	Form of a pause and a broken-off word: … or ... (three dots). Changing it does not correct existing text.
Autosave current work	on, every 60 s	A copy of your work in case of a crash (30–600 s, chapter 9).
Recording details	8 fields	Recording details fields, their order, personal fields and your team's own fields (chapter 7).
Models folder	application data folder	Where downloaded ASR models are stored (e.g. an external drive). Changing it does not move models already downloaded.
Device	automatic	An NVIDIA graphics card when available, otherwise the processor; or always the processor (chapter 8).
Vocabulary	empty	Proper names and dialect words for all recordings; passed to the model together with the project vocabulary (chapter 8).

Open the window with the sliders button or `Ctrl+,`. The settings are grouped into tabs: Appearance, Player, Notation, Project, Recording details and ASR. Changes take effect as soon as you click **OK** (except the language, which needs a restart) and are remembered after restarting. **Restore defaults** goes back to the values in the table.

The playback buttons are inactive until you load a recording. Loading a new recording while you have unsaved work asks whether to discard it.

Ellipsis style

The team writes a pause and a broken-off word with a single ellipsis character **…**. If you prefer three dots `...`, change it in the settings (Notation tab). The validator understands both forms, but shows ellipses written differently from the chosen form as hint VAL-05. The **Unify ellipses** button above the warning list replaces them all at once — a single `Ctrl+Z` undoes the whole replacement. This is particularly useful for texts written in an older version of the application. The note `[...]` and the marker `[…?]` have a fixed form and are not changed.

Find and replace

`Ctrl+F` opens the search bar below the editor; matches are highlighted (on screen only — the highlight does not go into the DOCX), and a counter shows which one is current. Options: **Aa** — match case, **Whole words**, **.*** — regular expression. An invalid expression does not damage the text: the bar shows a message.

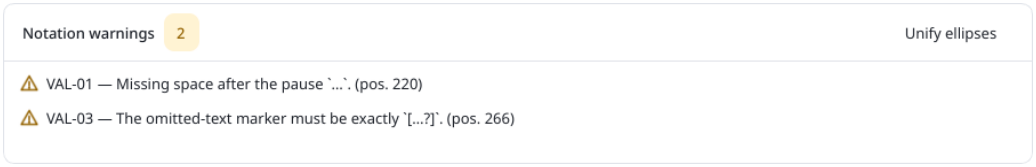
`Ctrl+H` adds the **Replace with** field. The ^ character makes the next letter superscript: replacing `będzie` with `be^ndzie` gives be^{n}dzie (a preview is shown next to the field). Type a literal ^ as \^. **Replace all** is a single undo step. The replacement is always literal, even with regular expressions.

A typical workflow

1. In the header of the **Transcription** card type the **Author**; check the **Date** (today by default, with a drop-down calendar).
2. Click **Import recording…** and choose the recording. The reference code in the recording details is suggested from the file name; fill in the remaining fields (chapter 7).
3. Save the project right away (`Ctrl+S`) — preferably next to the recording.
4. Play short fragments (`Ctrl+Space` or `F4`) and type them in the editor; `Alt+←` skips back a few seconds. If needed, slow down the speed or loop a difficult fragment.
5. Mark pronunciation phenomena: select nasal letters and an inserted *u* and press `Ctrl+Shift+↑` ; insert pauses and markers with the shortcuts.
6. Keep an eye on the **Notation warnings** card below the editor.
7. Save the project from time to time (`Ctrl+S`). The next day open it (`Ctrl+O`) — you will be back at the same position in the recording.
8. Finally: **Export DOCX…** (`Ctrl+E`).

Notation warnings

Below the editor the application continuously shows places where the notation may be inconsistent; the number next to the title says how many there are (green — no warnings). Click a warning to select the fragment in the text. Warnings never fix anything by themselves — if the notation is intentional, simply ignore them; they do not block the export. Drag the border between the editor and the warnings to enlarge the list.



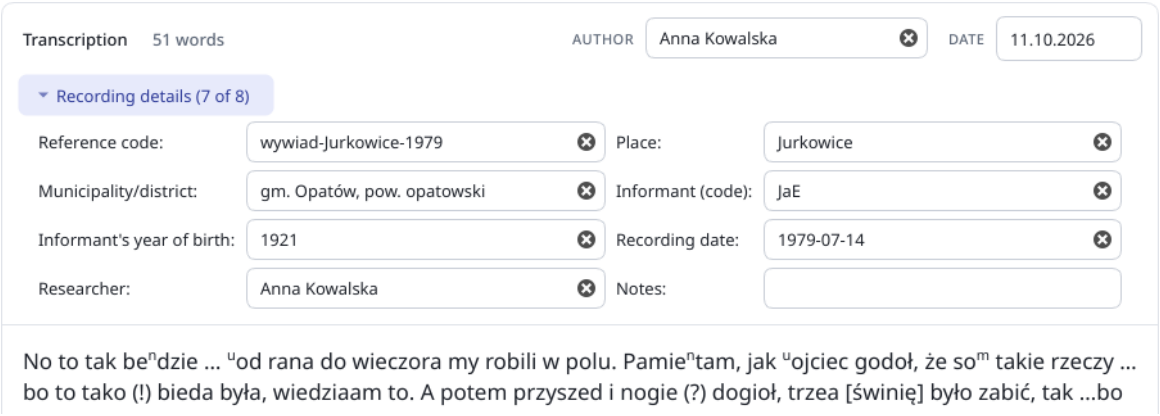
The Notation warnings card

Code	Meaning	How to fix
VAL-01	Pause without a space after the ellipsis: `tak …bo`	add a space: `tak … bo`
VAL-02	Note without a closing bracket: `[o żniwach`	close it: `[o żniwach]`
VAL-03	Malformed omission marker, e.g. `[…]`	use exactly `[…?]` (`Ctrl+3`)
VAL-05	Hint: an ellipsis written differently from the form chosen in the settings	**Unify ellipses** or fix it manually

An ellipsis attached to a word (`jeż…` `kam…` or `jeż...` `kam...`) marks a broken-off word and is not reported.

7. Recording details

The recording details are descriptive data about the recording: reference code, place, informant, recording date and researcher. They are saved in the project and go into a table at the beginning of the DOCX file. Expand them by clicking **Recording details** below the header of the Transcription card; next to it you can see how many fields are filled in. The transcription author and date fields stay on top.



Expanded recording details

Field (default)	Notes
Reference code	Suggested from the recording file name on import when the field is empty (`AdK_1954.aac` → AdK_1954). In the DOCX it becomes the document title.
Place, Municipality/district	The place also goes into the DOCX keywords.
Informant (code), Informant's year of birth	**Personal data** — left out of the anonymized export.
Recording date, Researcher, Notes	Free text.

Your team's own fields

In **Settings** (`Ctrl+,`) the **Recording details** tab shows the list of fields. Uncheck a field to hide it in the form and in the export; **Move up** / **Move down** change the order; the **Personal data** column decides about anonymization. **Add field…** creates your own field (e.g. tape number) — rename it with a double-click, and **Remove field** removes it from the list. Default fields cannot be removed, only hidden. Values of hidden fields are not lost — they stay in the project.

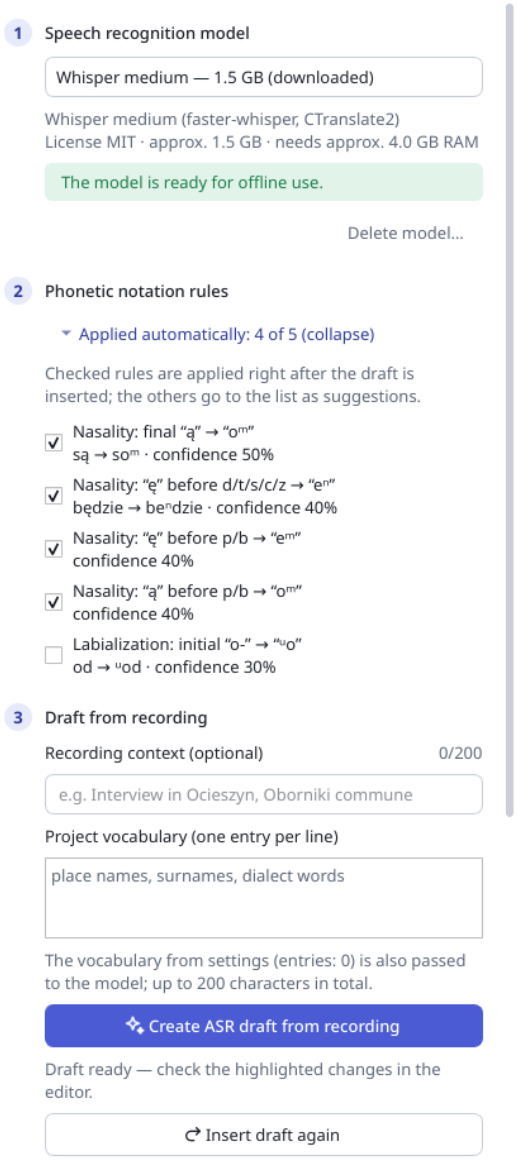
The names of the default fields follow the interface language; the names of your own fields stay as you typed them.
During ASR transcription the recording details, like the whole window, are locked.

8. ASR draft (automatic speech recognition)

The ASR draft is an **optional** helper: the Whisper speech recognition model listens to the recording and produces ordinary (orthographic) Polish text, which goes straight into the editor. Selected phonetic notation rules are applied automatically, and every change is highlighted for your review. Without the model the application works normally.

> **Remember.** The draft is an orthographic, not a phonetic, record. The model does not hear nasality, labialization or the reduction of “ł”, and dialect speech is recognized worse than standard Polish. The rules only guess — every fragment needs your correction.

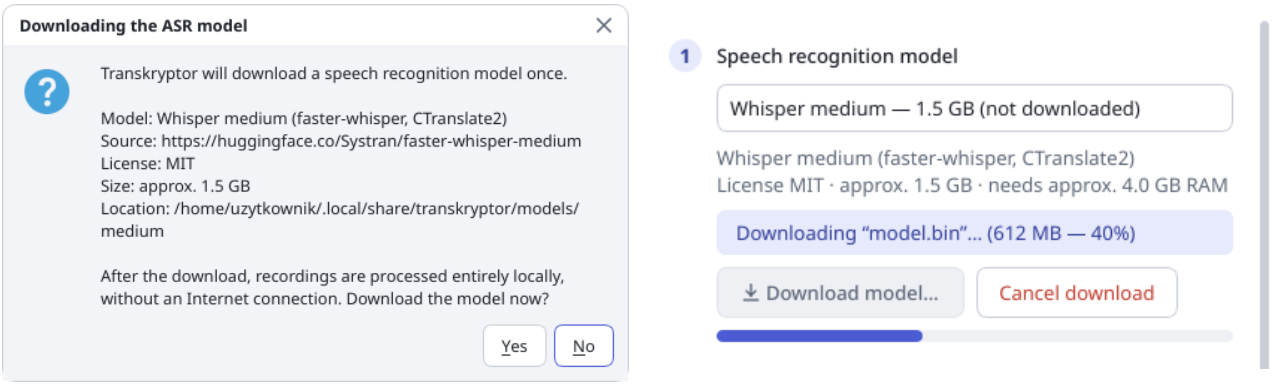
The panel guides you through three numbered steps: model, rules and draft. Below them are the list of segments and the list of changes to review.



Steps 1–3 in the ASR draft panel (rule list expanded)

Step 1 — choosing and downloading the model (once)

1. In the **ASR draft** panel choose a model from the list: **small** (approx. 0.5 GB, fast, for weaker computers), **medium** (approx. 1.5 GB, default) or **large-v3** (approx. 3.1 GB, most accurate, but slower and needs approx. 6 GB of RAM). The choice is remembered. Then click **Download model…**.
2. Read the information: model name, source, license (MIT), size and location. Click **Yes** to download or **No** to cancel.
3. Watch the progress bar. **Cancel download** stops the operation; after an error you can try again — an incomplete download is never treated as a ready model.
4. Once you see the green message **The model is ready for offline use.**, the Internet is no longer needed.



Download consent (left) and download in progress (right)

Each model is downloaded separately. Remove a model you no longer need with **Delete model…** — the application asks for confirmation and shows how much space was freed; other models keep working. If the selected model needs more memory than the computer has, a warning appears before transcription.

On a computer with an NVIDIA graphics card the draft can be created many times faster. This needs the CUDA 12 and cuDNN 9 libraries, installed by you (they are not bundled with the application). Without them the draft is computed on the processor, and a note appears below the transcription button. In **Settings** (**ASR** tab) you can force the processor only.

Step 2 — choosing the phonetic notation rules

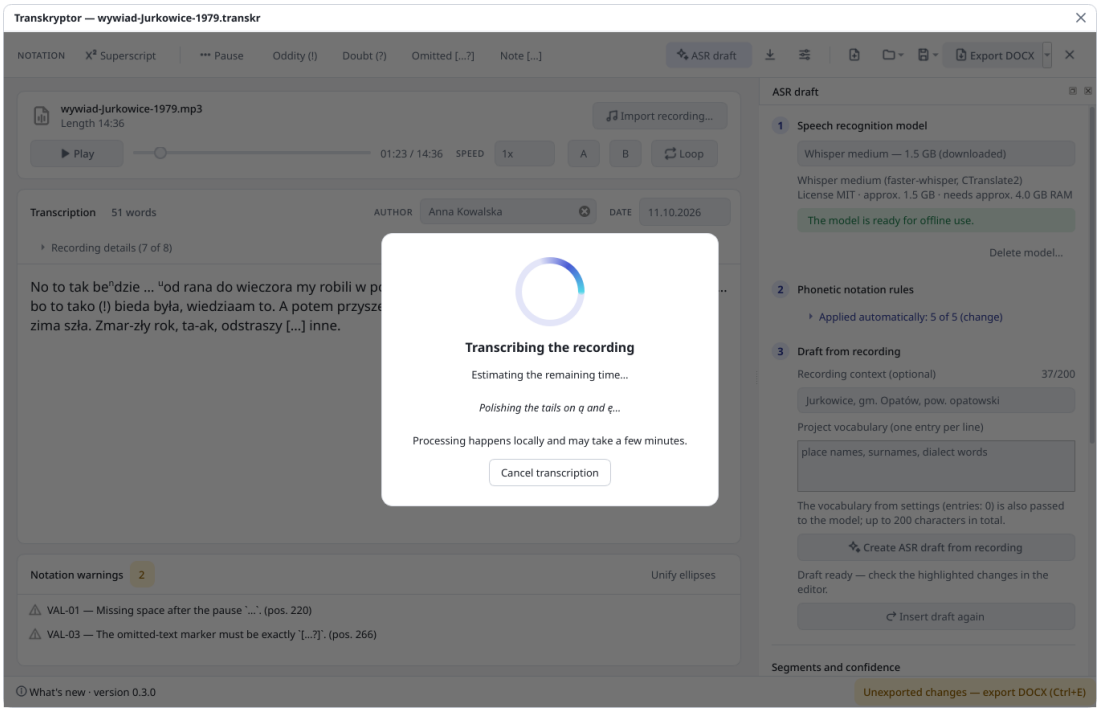
The rule list is collapsed by default — you only see how many rules are turned on. Click **Applied automatically: … (change)** to expand it. Each rule shows an approximate confidence. **Checked** rules are applied to the draft automatically; matches of **unchecked** rules appear in the list as suggestions for you to decide on. You can change the selection before every draft.

Step 3 — creating the draft

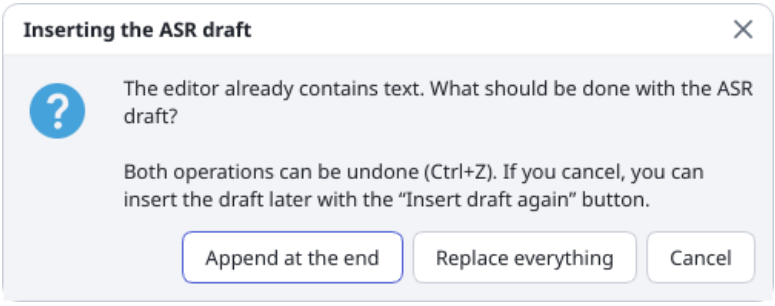
1. Import the recording (**Import recording…** button in the player).
2. Optionally fill in **Recording context** (up to 200 characters; filled in automatically from the place and commune in the recording details) and **Project vocabulary** — place names, surnames, dialect words, one per line. They help the model spell proper names; the project vocabulary is saved in the project file.
3. Click **Create ASR draft from recording**. The button is active when the model is downloaded and a recording is loaded.

4. During the transcription the window is locked and a progress card appears in the middle. Stop it with **Cancel transcription** — the current fragment will be finished first.
5. When it is done, the draft goes into the editor. If the editor already contained text, the application asks whether to **Replace everything** or **Append at the end**; **Cancel** leaves the text unchanged.

Careful with hints. A short, factual context helps, but a long or misleading one may make the model repeat it in the draft. If text from the context that is not in the recording shows up in the draft, shorten or clear the field.



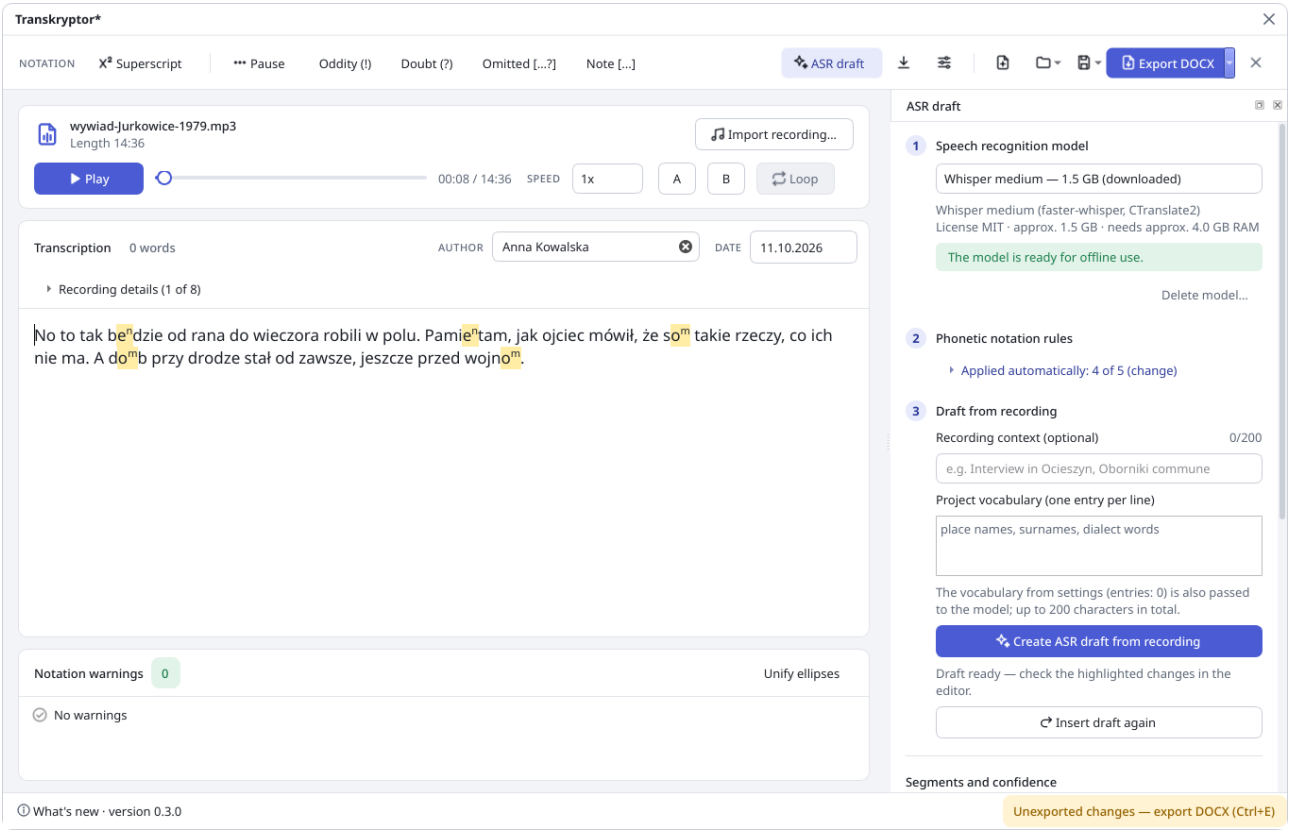
The window during transcription — only cancelling is active



Question about where to put the draft when the editor contains text

Reviewing the changes

Changes made by the rules are highlighted in yellow in the editor. The **Changes and suggestions to review** list shows every match: the word before and after the change, the kind of phenomenon and the confidence.



Draft inserted into the editor — highlighted changes and the review list

Element	Action
✓ (green)	Change applied automatically — highlighted in the editor.
○	Suggestion of an unchecked rule — the text has not been changed.
Clicking an item	Selects the fragment in the editor so that you can easily listen to it and correct it.
Double-click / Apply suggestion	Applies the ○ suggestion to its fragment. If you have changed the fragment since the transcription, the application changes nothing and asks you to correct it manually.
Finish review	Removes the highlights and the list; the text stays unchanged.
Insert draft again	Inserts the last draft once more without repeating the transcription — also after reopening the project.
Segments and confidence	Draft fragments with their time (mm:ss–mm:ss) and approximate confidence. Low-confidence segments are highlighted in colour — listen to them especially carefully. Clicking a segment (or Enter) plays the recording from its start. After importing a different recording the list is cleared.

> **Undo.** Inserting the draft and applying the rules are two separate steps: the first `Ctrl+Z` restores the raw draft without the rules, the second — the state before the insertion.

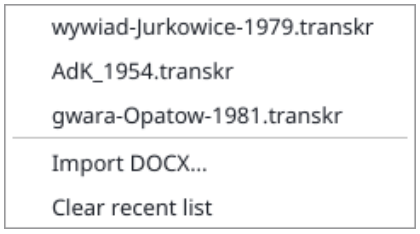
An unfinished review, the segment list and the last draft are saved in the project: when you open it, the highlights and the list come back and you can keep applying suggestions.

9. Project and autosave

The project file (**.transkr**) keeps all the work on a recording: the text with superscripts, the author, date and recording details, a reference to the recording, the playback position and A–B loop, the last ASR draft with its segments and an unfinished rule review. The recording itself is not copied into the project — the project remembers where it is.

Saving and opening

Action	How
Save project	`Ctrl+S` or the floppy disk button. The first save asks for a name and folder (the suggested name is the reference code).
Save project as…	`Ctrl+Shift+S` or the arrow next to the floppy disk — a copy in a new file; you continue working in the new file.
Open project…	`Ctrl+O` or the folder button.
Recent projects	The arrow next to the folder — up to 8 recently saved or opened projects.



The Open project button menu

> **Keep the project next to the recording.** The project remembers where the recording is relative to itself. When you move the whole folder (also to another computer, or from Windows to Linux), the project still finds the recording.

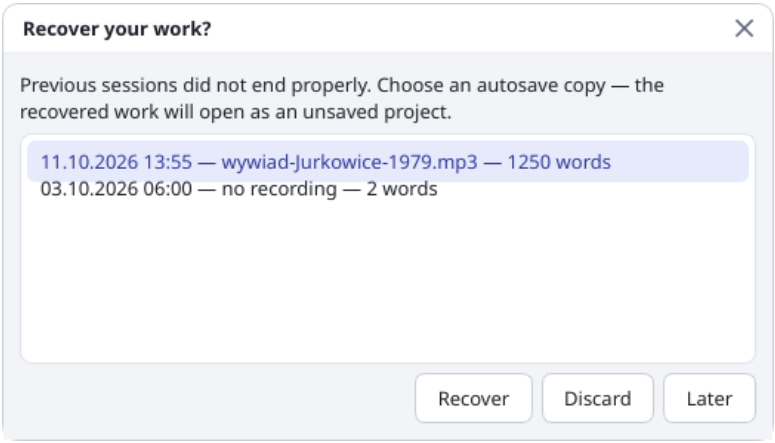
When the recording is missing or different

- **Recording moved or deleted** — the application asks whether to locate the file. **Yes** opens a file dialog, **No** opens the project without the recording: you can keep editing the text and import the recording later.
- **Recording changed or replaced** — the application compares the file checksum. If it does not match, it warns that the segment times and the position may not fit, and asks whether to use the recording anyway.

Autosave

Every 60 seconds (change the interval in the settings, Project tab) the application saves a copy of your current work in its data folder — only when something has changed. Autosave does not change your project file and does not remove the asterisk from the title. During ASR transcription it waits until the end. After a clean exit the copy is deleted.

If the application ends abnormally (a freeze, a power cut), the next time you start it you will see a window with a list of copies: date, recording and number of words.



The recovery window after a crash

Button	Effect
Recover	Opens the selected copy as a new, unsaved project. Save it right away (`Ctrl+S`) — e.g. in place of the previous project.
Discard	Deletes the selected copy permanently.
Later	Keeps the copies — the window appears at the next start.

10. Export to DOCX

1. Click **Export DOCX…** or press `Ctrl+E` .
2. Choose the folder and the file name (suggested: the reference code from the recording details or `transcription-YYYY-MM-DD.docx`). The .docx extension is added automatically.
3. Click **Save**. The status bar shows “No unexported changes”.

What the file contains

- a recording details table — only the filled-in fields, in the order from the settings,
- an *Author: … Date: …* paragraph (if they are filled in),
- the transcription text with all Polish characters, hyphens and markers,
- superscripts as Word formatting (superscript), not as separate characters,
- in the document properties: the author, the title (the reference code or “Phonetic transcription”) and keywords (the place).

The table labels, the *Author/Date* paragraph and the default title are written in the interface language. Transkryptor imports documents exported in either language.

The ASR review highlights do not go into the file. Open the file in Microsoft Word, LibreOffice Writer or Google Docs. If saving fails (e.g. no permission for the folder), you will see a message with a hint — choose a different location and try again.

Anonymized export

The arrow next to the **Export DOCX…** button → **Export anonymized DOCX…** creates a file without the recording details fields marked as personal data (by default the informant's code and year of birth) — also in the document properties. The text and superscripts are the same as in a normal export. The suggested file name ends with `-anonymized.docx`. Choose which fields are personal in the Recording details tab of the settings.

> **Note.** Anonymization only applies to the recording details. You have to remove names of people or places from the transcription text itself manually.

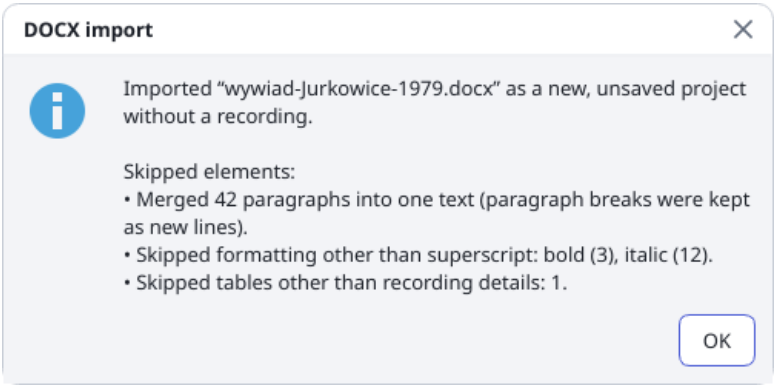
11. DOCX import

You can return to a transcription exported to DOCX earlier (or typed in Word) by importing it: the arrow next to the **Open project** button → **Import DOCX…**. The application creates a new, unsaved project without a recording — import the recording and save the project. If you have unsaved work, it asks first, as with a new document.

What is loaded

- the text and superscripts,
- the recording details from the table at the beginning of the document (recognized by field names in English or Polish, including your team's own fields),
- the author and date from the *Author: … Date: …* paragraph (or its Polish version *Autor: … Data: …*); if it is missing — the author from the document properties.

A file exported from Transkryptor comes back complete. Other documents are loaded on a best-effort basis: several paragraphs are joined into one text (paragraph breaks stay as new lines), and bold, italics, underline, tables other than the recording details, images and footnotes are skipped. After the import the application shows what it skipped.

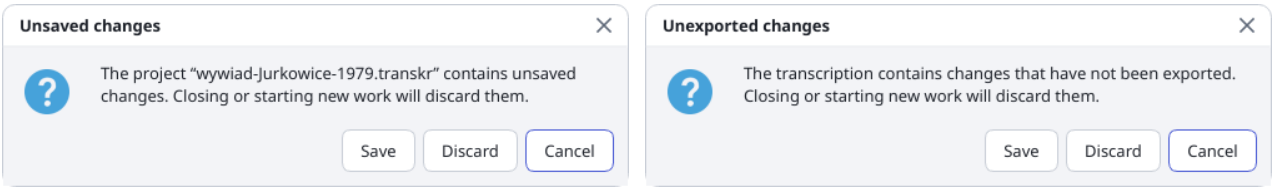


Report after importing a document

12. Closing and unsaved changes

When you close the application, create a new document, load another recording, open a project or import a DOCX, the application checks that no work will be lost:

- **The work has a project file** — it only asks about changes not saved in the project. A saved project that has not been exported to DOCX closes without asking.
- **The work has no project file** — it asks about changes not exported to DOCX.



The question for a project (left) and for work without a project file (right)

Button	Effect
Save	Saves the project (without a file — asks for a name first) and then continues the operation.
Cancel	You go back to work — nothing is lost.
Discard	The changes are discarded permanently and the operation continues.

The application cannot be closed during ASR transcription — cancel the transcription first or wait for it to finish.

13. Privacy

- Recordings and text never leave your computer. There is no cloud speech recognition and no synchronization.
- The application does not collect statistics (telemetry).
- The application connects to the Internet for two purposes only: downloading the ASR model (only after you confirm) and checking whether a newer version is available (at most twice a day, without sending any text, recordings or file paths).
- After the model has been downloaded, everything works offline too.
- You save projects and DOCX files yourself, wherever you choose.
- **Autosave** creates a copy of your work on disk — the text, the recording details (including the informant's code) and the recording path — in the application data folder (table below, `autosave` subfolder). The copy is removed after a clean exit, a new document and saving the project. You can turn autosave off in the settings (Project tab).
- Settings (player, font, language, recording details fields, recent projects list) are stored on your computer: Linux — `~/.config/transkryptor`, Windows — the registry key `HKCU\Software\transkryptor`. They contain no text and no recording details values.

Responsibility for the consent of the recorded people and for data protection rules lies with the person or institution doing the transcription.

Where the ASR model and autosave copies are

System	Application data folder
Windows	`%APPDATA%\Transkryptor`
Linux	`~/.local/share/transkryptor`

The ASR model is in the `models` subfolder, autosave copies — in `autosave`.

Remove a single model in the **ASR draft** panel with **Delete model…**. Deleting the whole `models` folder brings back purely manual work and does not affect anything else. The models folder can be moved in the settings (**ASR** tab).

14. Troubleshooting

Problem	What to do
I cannot choose the recording file	The application accepts MP3, WAV, FLAC, M4A, AAC, OGG, OPUS and WMA. Convert other formats, e.g. in Audacity or VLC.
There is no sound	Check the system volume and the output device. On Linux — whether PipeWire/PulseAudio is running.
“Create ASR draft” is greyed out	Download the model and import a recording first.
Downloading the model failed	Check your Internet connection and that you have 3 GB of free space, then click **Download model…** again.
I cannot type during transcription	This is intentional — the window is locked until the transcription ends. Wait or click Cancel transcription.
Transcription takes a very long time	This is normal on a slower processor. You can cancel and try a shorter recording.
A suggestion was not applied	The fragment has changed since the transcription — correct it manually (click the item to select it).
The ASR draft panel disappeared	Click **ASR draft** in the toolbar.
Superscript does not work	Select the letters before pressing `Ctrl+Shift+↑`.
The interface is in the wrong language	**Settings → Appearance → Language**, then restart the application.
Linux: the AppImage does not start	Make it executable (chmod +x), install libfuse2. glibc 2.35+ is required (Ubuntu 22.04 or later).
Windows: SmartScreen warning	More info → Run anyway.
`Ctrl+Space` does not start playback	Use `F4` — on Linux Ctrl+Space is sometimes taken by the input method. The player shortcuts also do not work during ASR transcription.
The project opened without the recording	The recording was moved and “No” was chosen when opening. Click **Import recording…** and save the project.

Problem	What to do
“Different recording” warning	The recording file differs from the one saved in the project (e.g. it was trimmed or converted). Use it if you are sure it is the same recording.
The application froze or shut down	Start it again and choose Recover (chapter 9).
Bold disappeared after a DOCX import	A transcription only keeps superscript — other formatting is skipped (chapter 11).
Is the application working correctly?	Run it with --self-test (chapter 3).

15. Uninstalling

- **Windows:** Settings → Apps → Installed apps → Transkryptor → Uninstall (this also removes the settings). Then, if you want, delete the `%APPDATA%\Transkryptor` folder (ASR model, autosave copies).
- **Linux:** delete the AppImage file and, optionally, the `~/.local/share/transkryptor` (ASR model, autosave copies) and `~/.config/transkryptor` (settings) folders.

Your projects and DOCX files are not deleted.

Appendix: notation cheat sheet

The phenomena below belong to the Polish dialectological notation used by the application; the examples are Polish words.

Phenomenon	Notation	How in the application
Nasality	be^{n}dzie, pamie^{n}tam, so^{m}	select the letter → `Ctrl+Shift+↑`
Diphthongization, labialization	^{u}od, sw^{u}oje, k^{u}osa	select *u* → `Ctrl+Shift+↑`
Reduction of “ł”	wiedziaam, wyczytaam	write without “ł” — the application does not correct it
Stress group	zmar-zły, ta-ak	a plain hyphen
Modified close vowel	uóna, sóm, słóńce	the letter ó
Pause	rzeczy ... bo	Pause / `Ctrl+.`
Broken-off word	jeż... kam...	ellipsis with no space before it
Pronunciation oddity	zamudze (!)	Oddity / `Ctrl+1`
Uncertain hearing	nogie (?)	Doubt / `Ctrl+2`
Omitted text	odstraszy [...?] inne	Omitted / `Ctrl+3`
Transcriber's note	trzea [świnię]	Note / `Ctrl+[`

Transkryptor is open-source software licensed under GPL-3.0-or-later. The licenses of the libraries and the model are in the `licenses` folder next to the application.