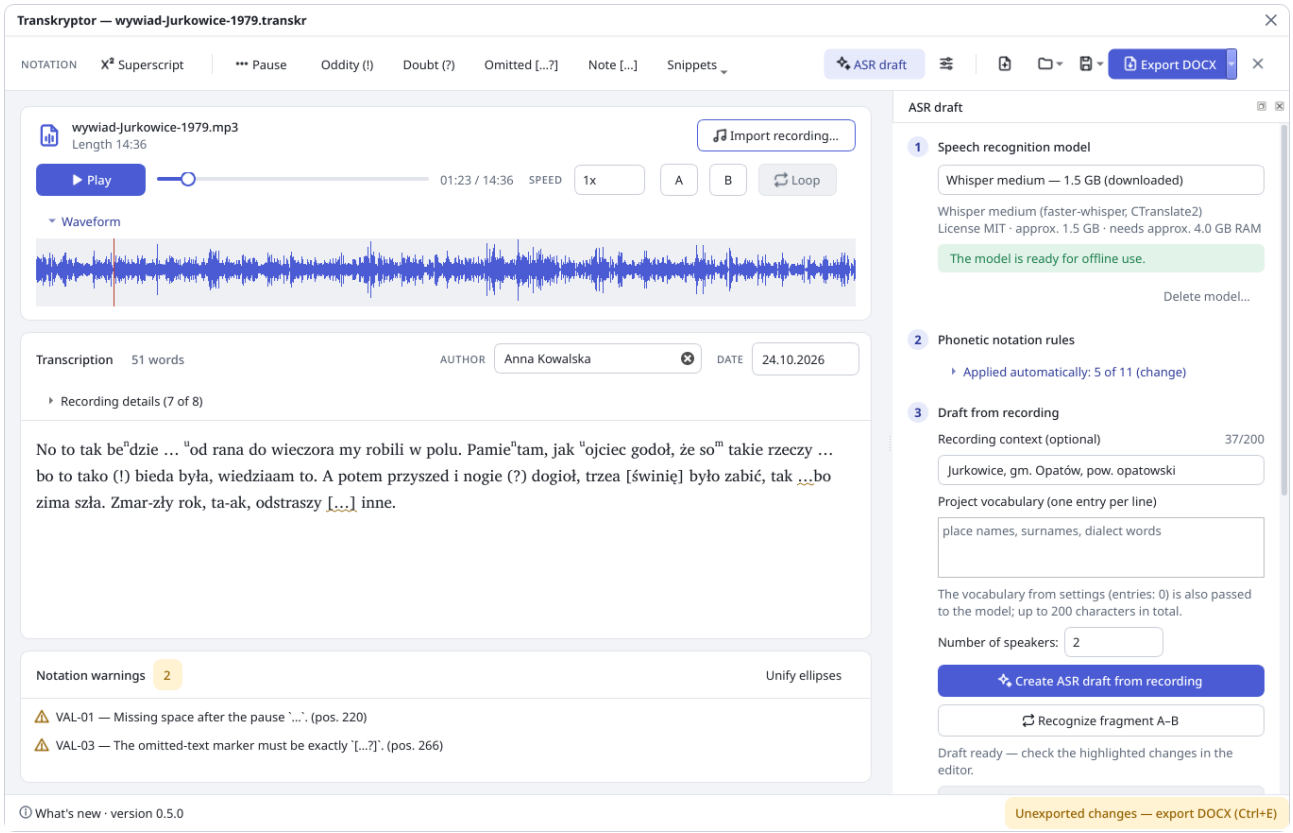
TRANSKRYPTOR

User guide

A local desktop application for manual phonetic transcription of Polish speech, assisted by an optional automatic speech recognition draft.



Transkryptor at work

Application version 0.5.0 · Windows 10/11 (64-bit) and Linux x86_64 · October 2026

Contents

1. What Transkryptor is . . . 3
2. Requirements . . . 3
3. Installation . . . 4
4. The application window . . . 5
5. Buttons and keyboard shortcuts . . . 7
6. Working with the recording and the text . . . 9
7. Recording details . . . 14
8. ASR draft (automatic speech recognition) . . . 16
9. Tailoring the notation: validator, rules, shortcuts . . . 22
10. Project and autosave . . . 25
11. Export: DOCX and other formats . . . 27
12. DOCX import . . . 30
13. Closing and unsaved changes . . . 30
14. Privacy . . . 31
15. Troubleshooting . . . 32
16. Uninstalling . . . 33
Appendix: notation cheat sheet . . . 34

1. What Transkryptor is

Transkryptor helps you write down a recorded utterance the way it sounds, not the way it is spelled. You listen to an MP3 recording, type the transcription in the editor, mark pronunciation phenomena (nasality, an inserted *u*, pauses, uncertain passages), and export the finished text to a Word file (DOCX).

The application runs entirely on your computer. There is no account, no sign-in, and recordings are never sent to the Internet.

Transkryptor is built for Polish dialect speech, so the notation and the examples in this guide are Polish. The user interface is available in English and Polish — see **Settings → Appearance → Language** (chapter 6).

What it can do

- play an MP3 recording with adjustable speed (0.5×–2×) and seeking,
- insert notation markers with a single shortcut: pause, (!), (?), [...?], [...],
- mark letters as superscript (e.g. be^{n}dzie, ^{u}od),
- warn about notation errors as you type (e.g. an unclosed bracket),
- export the transcription together with the author, date and recording details to DOCX, keeping superscripts — also in an anonymized version,
- save your work in a project file and come back to it later; autosave protects it against crashes,
- load a previously exported (or someone else's) DOCX document as a new project,
- optionally: create an automatic text draft (ASR), insert it straight into the editor and apply selected phonetic notation rules — with every change highlighted for your review,
- follow the light or dark theme of your system.

> **Important.** You save your work in a project file (**Save project**, `Ctrl+S`) — with the text, recording details, ASR draft and the position in the recording. The DOCX file is the result for other people. Autosave is only a safety net in case of a crash; it does not replace saving the project (chapter 10).

2. Requirements

	Minimum	Recommended
System	Windows 10/11 64-bit or Linux x86_64 (glibc 2.35+, e.g. Ubuntu 22.04 or later)	
RAM	4 GB	8 GB
Disk space	approx. 0.5 GB (application only); 3 GB with the ASR model	
Processor	4 cores	8 cores or more
Internet	only once, to download the ASR model (1.5 GB)	

You do not need Python or any other tools — the installer contains everything. On Linux, playback needs working system audio (PipeWire or PulseAudio; typical distributions have it by default).

Manual work (typing, listening, export) runs on much weaker hardware. The requirements mainly concern the ASR draft. On the test machine an hour of recording was processed in about half an hour; a slower processor will take longer.

3. Installation

Download the file for your system and the **SHA256SUMS.txt** file.

System	File
Windows	Transkryptor-0.5.0-windows-x64-setup.exe
Linux	Transkryptor-0.5.0-x86_64.AppImage

Windows

1. Double-click **Transkryptor-0.5.0-windows-x64-setup.exe**.
2. If a **Windows protected your PC** window (SmartScreen) appears, click **More info** → **Run anyway**. The application is not signed with a commercial certificate, hence the warning.
3. By default the installer installs the application for your account only and does not need administrator rights. In the first window you can choose to install it for all users instead.
4. Accept the license (GPL-3.0), check the desktop shortcut option if you want one, and click **Install**.
5. At the end you can start the application right away. Later you will find it in **Start menu** → **Transkryptor**.

Linux (AppImage)

An AppImage is a single file that needs neither installation nor administrator rights.

1. Save the file, e.g. in the `~/Applications` folder.
2. Make it executable — in the file manager: **Properties** → **Permissions** → **Allow executing**, or in a terminal:

```
chmod +x Transkryptor-0.5.0-x86_64.AppImage
./Transkryptor-0.5.0-x86_64.AppImage
```

- If the application does not start and a message about FUSE appears, install the `libfuse2` package (Ubuntu: `sudo apt install libfuse2`).
- You can add a system menu entry with a tool such as Gear Lever or AppImageLauncher.

Verifying the downloaded file (optional)

The checksum confirms that the file has not been damaged or tampered with.

```
# Linux
sha256sum --check SHA256SUMS.txt

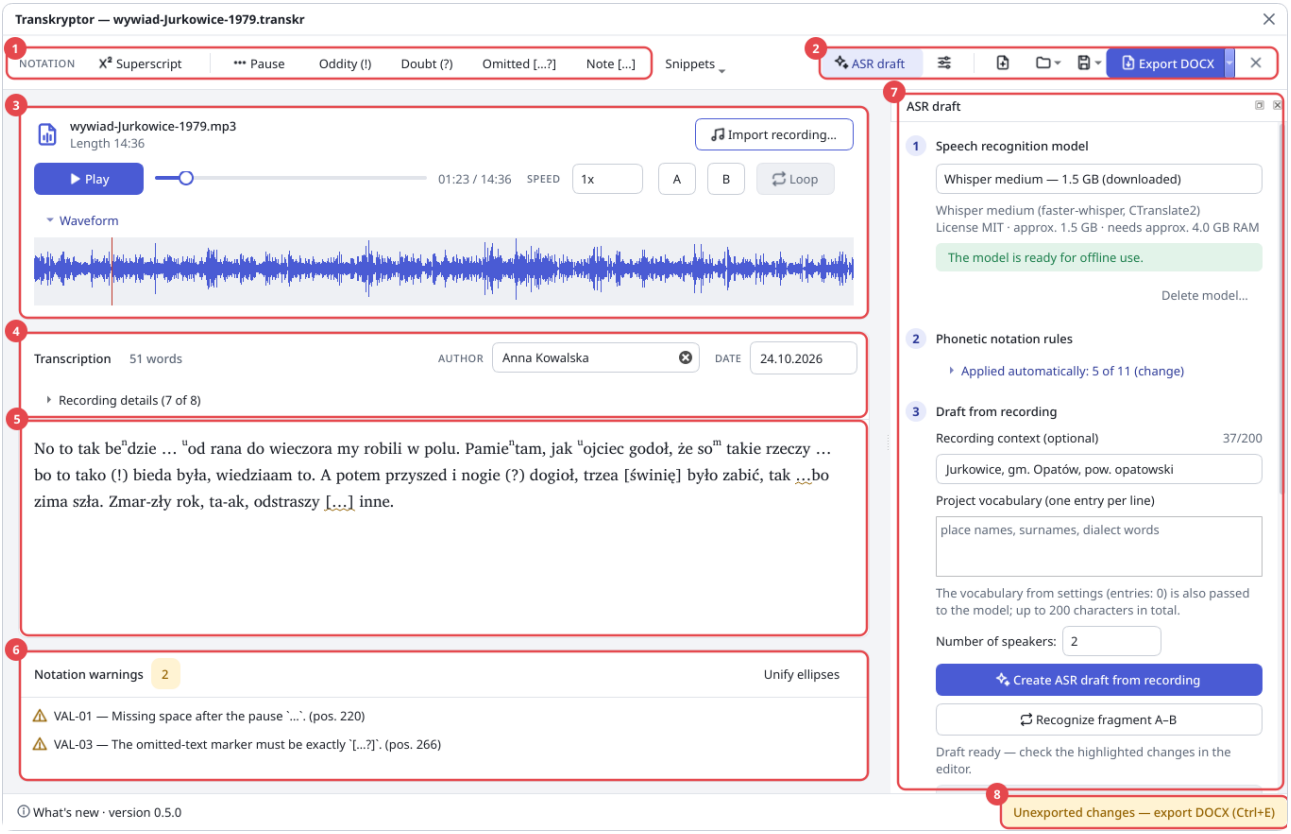
# Windows (PowerShell) — compare the result with SHA256SUMS.txt
```

```
Get-FileHash .\Transkryptor-0.5.0-windows-x64-setup.exe -Algorithm SHA256
```

Diagnostics without opening the window: run the application with the `--self-test` switch. Every item should have the `[ok]` status; the line “ASR model: not downloaded” is not an error.

4. The application window

After starting the application you will see one window: the toolbar at the top, the working area on the left (player, transcription, warnings) and the **ASR draft** panel on the right. Drag the borders between the areas with the mouse to change their proportions.



The main window at work — the numbers match the table below

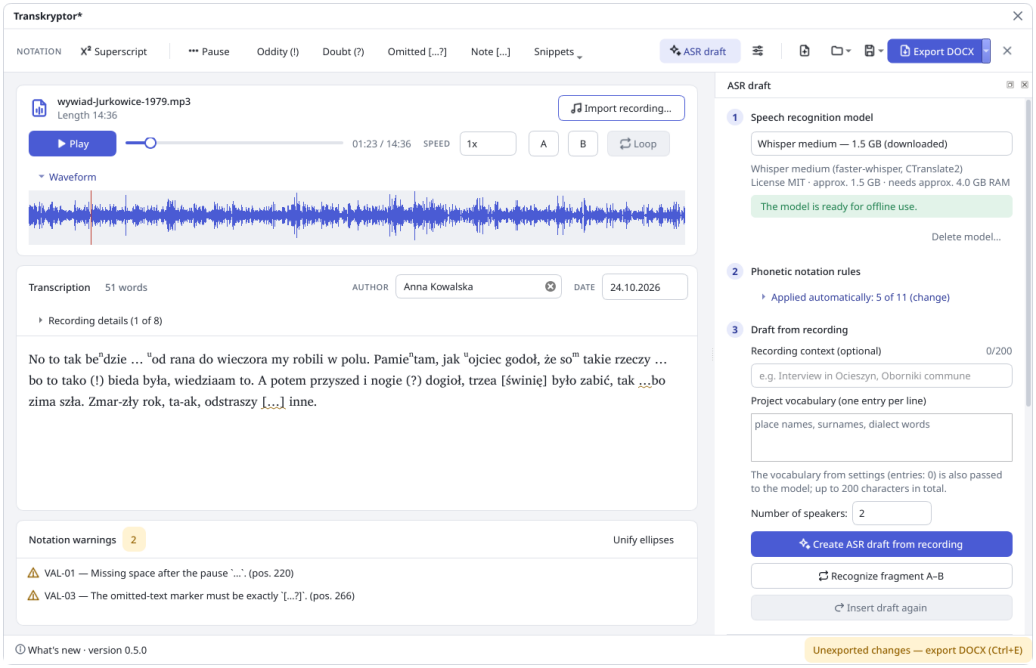
No.	Element	What it is for
1	Toolbar — Notation	superscript and inserting markers (chapter 5)
2	Toolbar — actions	ASR panel, updates, settings, new document, opening and saving the project, export to DOCX (highlighted button), closing the application
3	Player	recording import, name and length, playback, seeking, speed
4	Transcription header	word counter, the Author and Date fields and the collapsible recording details (chapter 7) — saved in the project and in the DOCX file
5	Editor	the place for the transcription text

No.	Element	What it is for
6	Notation warnings	list of detected deviations from the notation, with a counter
7	ASR draft panel	model download, rules, automatic draft, change review (chapter 8)
8	Status bar	on the left “What's new” (the change history), on the right — whether all changes have been exported

> **Tip.** The window title shows the name of the open project. An asterisk (e.g. *Transkryptor — interview.transkr**) means unsaved project changes or, when the work has no project file yet, unexported changes. A yellow message in the status bar reminds you to export to DOCX.

Light and dark theme

By default Transkryptor follows the system theme: when dark mode is selected in Windows or your Linux desktop, the application window is dark too. You can also choose the light or dark theme permanently in **Settings → Appearance**. The look does not affect the exported DOCX file.



The same window in the dark theme

The ASR draft panel

You can detach the panel into a separate window or close it with the buttons in its corner. Bring a closed panel back with the **ASR draft** button in the toolbar.

5. Buttons and keyboard shortcuts

Notation



Toolbar — the Notation group

Button	Shortcut	What it does	Example
Superscript	`Ctrl+Shift+↑`	Turns the selected letters into superscript (or back). Nothing happens without a selection.	be^{n}dzie, ^{u}od
Pause	`Ctrl+.`	Inserts a pause: an ellipsis with a space on both sides, in the form chosen in the settings (… by default).	takich rzeczy … bo
Oddity (!)	`Ctrl+1`	Unusual pronunciation — the spelling is not a mistake.	zamudze (!) pani
Doubt (?)	`Ctrl+2`	Uncertain hearing.	nogie (?) dogioł
Omitted […?]	`Ctrl+3`	Inserts […?] — a fragment left out because it was unintelligible.	odstraszy […?] inne
Note […]	`Ctrl+[`	Inserts […] and puts the cursor inside — type your comment.	trzea [świnię]

The tables in this chapter list the default shortcuts. Every shortcut can be changed in **Settings…** (**Shortcuts** tab), and frequently used forms can be inserted with the **Snippets** button or your own shortcuts — see chapter 9.

Actions



Toolbar — the actions on the right

Button	Shortcut	What it does
ASR draft	—	Shows or hides the ASR draft panel.
Down arrow (updates)	—	Checks whether a newer version of the application is available.
Settings (sliders)	`Ctrl+,`	Opens the settings for appearance and language, player, notation, project and recording details (chapter 6).
New document (sheet with a plus)	`Ctrl+N`	Clears the editor and the metadata and closes the recording. If you have unsaved work, it asks first.
Open project (folder)	`Ctrl+O`	Opens a .transkr project file. The arrow next to it opens a menu: recent projects, Import DOCX… (chapter 12), Clear recent list.

Button	Shortcut	What it does
Save project (floppy disk)	`Ctrl+S` `Ctrl+Shift+S`	Saves the project; the first time it asks for a file name. The arrow next to it: Save project as… (chapter 10).
Export DOCX…	`Ctrl+E`	Saves the transcription to a Word file. The arrow next to it: Export anonymized DOCX… (chapter 11).
× (Close)	`Ctrl+W`	Closes the application (with a warning about unsaved work).

Player

The player shortcuts work while the cursor is in the editor — you can control the recording without taking your hands off the keyboard. They do not insert characters into the text.

Shortcut	Action
`Ctrl+Space` or `F4`	play / pause (after a pause the recording rewinds by 1.5 s — auto-rewind)
`Alt+←` / `Alt+→`	skip the recording back / forward by 3 s (length in the settings)
`Ctrl+Shift+,` / `Ctrl+Shift+.`	slower / faster (next speed from the list)
`Ctrl+Shift+A` / `Ctrl+Shift+B`	loop start (A) / end (B) at the current position
`Ctrl+Shift+L`	turn the A–B loop on / turn it off and clear the markers
`Ctrl` + click in the text	play the recording from that place in the text (an ASR draft word — in a loop, until you pause)
`Ctrl+Shift+Space`	play the recording from the cursor position in the text
`Ctrl+Shift+T`	insert a timestamp of the current playback position at the cursor

Ctrl+Space on Linux. Some input methods (e.g. IBus, fcitx) capture `Ctrl+Space`. Use `F4` instead.

Find and replace

Shortcut	Action
`Ctrl+F`	opens the search bar below the editor
`Ctrl+H`	opens the bar with the “Replace with” field
`F3` / `Shift+F3`	next / previous match
`Esc`	closes the bar and returns to the editor

Speakers

Shortcut	Action
`Alt+Shift+1` … `Alt+Shift+9`	assign the speaker with that number in the list to the current paragraph (or the selected ones); again — remove it
`Alt+Shift+0`	remove the paragraph speaker

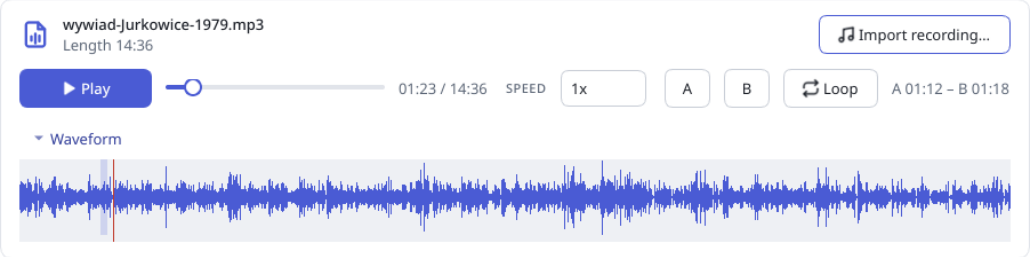
Standard editor shortcuts

Shortcut	Action
`Ctrl+Z` / `Ctrl+Shift+Z`	undo / redo — works for every change, including an inserted ASR draft and applied rules
`Ctrl+C` , `Ctrl+X` , `Ctrl+V`	copy, cut, paste
`Ctrl+A`	select all text

Tip. Hover the mouse over any button to see its description and shortcut. The marker shortcuts work while the cursor is in the editor — you do not need to take your hands off the keyboard while listening.

6. Working with the recording and the text

Player



The player card with a recording loaded

Element	Action
Import recording…	Opens a file dialog (MP3, WAV, FLAC, M4A, AAC, OGG, OPUS, WMA). The file is not copied — the application reads it from where it is.
Recording name	Name of the loaded file and length of the recording. Hover over it to see the full path.
Play / Pause	Starts or pauses playback (the label and icon change with the state).
Slider	Shows the position. Drag it to seek.
Time	Current position / length of the recording (mm:ss).
Speed	0.5x, 0.75x, 1x, 1.25x, 1.5x, 2x — a slower speed makes difficult passages easier to hear.

Element	Action
A, B	Mark the start and end of a fragment to loop at the current position. The range is shown next to them, e.g. A 01:12 – B 01:18.
Loop	Plays the A–B fragment over and over (without B — to the end of the recording). Clicking again turns the loop off and clears the markers.
Waveform	The shape of the recording — silence, pauses and loud passages are visible. A click seeks, dragging with the mouse sets the A–B loop, `Ctrl` + mouse wheel zooms (the wheel and the bar under the waveform scroll). ASR draft segments are shown at the bottom (green — confident, orange — uncertain), manual timestamps at the top. The arrow next to the title collapses the waveform.

The waveform is computed in the background after the recording is loaded (an hour of audio takes a dozen or so seconds; you can keep working) and is remembered, so it appears immediately the next time you open that recording.

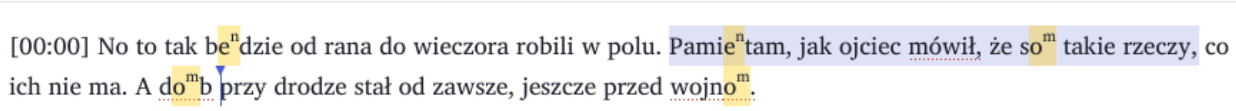
Loop and auto-rewind

The easiest way to listen to an unclear passage several times is a loop: at its start press `Ctrl+Shift+A`, at its end `Ctrl+Shift+B`, and then `Ctrl+Shift+L`. When playback reaches B it jumps back to A until you turn the loop off.

When you pause playback and resume it, the recording rewinds by 1.5 s so that you hear the context of the last words. Auto-rewind does not apply to the very first start or after seeking during a pause. You can change the rewind length in the settings or turn it off completely.

Text linked to the recording

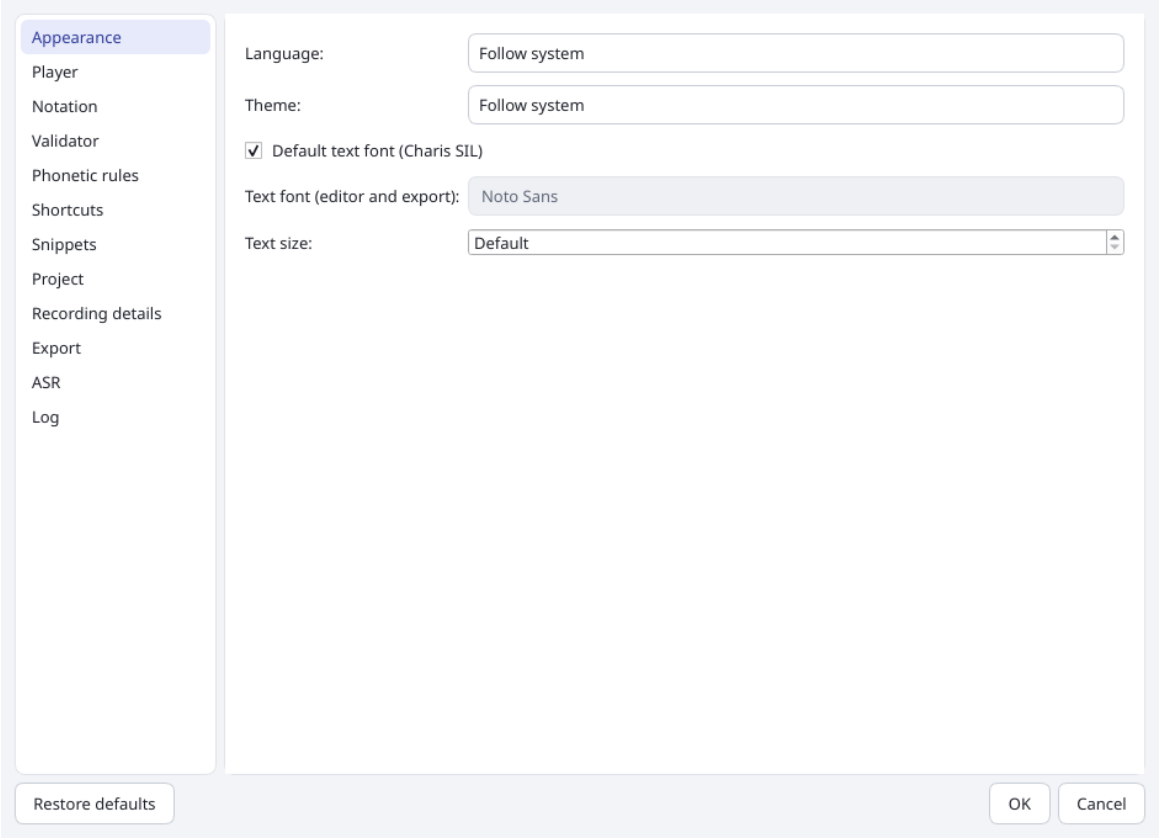
The ASR draft remembers where in the recording each sentence comes from. `Ctrl` + click in the text or `Ctrl+Shift+Space` plays the recording from the start of the sentence (0.5 s earlier, as set in the settings). During playback the current sentence is highlighted and the editor scrolls along. The links move with your corrections and are saved in the project.



The sentence being played (highlight), uncertain ASR words (dots) and a manual timestamp (triangle before “przy drodze”)

Text you type by hand can be linked to the recording with a timestamp: while listening, press `Ctrl+Shift+T` — a small triangle with the current playback position appears at the cursor. `Ctrl` + click next to the timestamp plays the recording from exactly that time. Remove a timestamp from the context menu (right mouse button). Timestamps and highlights are not included in the DOCX export.

Settings



The Settings window

Setting	Default	Meaning
Language	follow system	User interface language: English or Polish (Polski). The change takes effect after you restart the application. The DOCX export uses the interface language.
Theme	follow system	Light, dark or the same as the system.
Skip forward and back	3 s	Length of the `Alt+←` / `Alt+→` skip.
Auto-rewind	on, 1.5 s	How far the recording rewinds when resuming after a pause.
Start before ASR segment	0.5 s	A clicked draft segment starts slightly earlier so that the first syllable is not cut off.
Scroll the text along with playback	on	During playback the editor shows the draft sentence being played.
Editor font and text size	default	Look of the text in the editor. Does not affect the DOCX export.
Ellipsis style	… (single character)	Form of a pause and a broken-off word: … or ... (three dots). Changing it does not correct existing text.
Autocomplete	off	Suggestions of snippets and words from the document (chapter 9).

Setting	Default	Meaning
Validator	all rules except VAL-11	Which notation warnings are reported (chapter 9).
Phonetic rules, Shortcuts, Snippets	—	Your own notation rules, keyboard shortcuts and snippets (chapter 9).
Autosave current work	on, every 60 s	A copy of your work in case of a crash (30–600 s, chapter 10).
Recording details	8 fields	Recording details fields, their order, personal fields and your team's own fields (chapter 7).
Models folder	application data folder	Where downloaded ASR models are stored (e.g. an external drive). Changing it does not move models already downloaded.
Device	automatic	An NVIDIA graphics card when available, otherwise the processor; or always the processor (chapter 8).
Vocabulary	empty	Proper names and dialect words for all recordings; passed to the model together with the project vocabulary (chapter 8).
Word timing and confidence	on	Lets you play a single word and underlines uncertain words; transcription takes a few percent longer.
Uncertain words, threshold	underlined, 40%	Words with confidence below the threshold get a dotted underline (chapter 8). The threshold is approximate.

Open the window with the sliders button or `Ctrl+,`. The settings are grouped into tabs: Appearance, Player, Notation, Project, Recording details and ASR. Changes take effect as soon as you click **OK** (except the language, which needs a restart) and are remembered after restarting. **Restore defaults** goes back to the values in the table.

The playback buttons are inactive until you load a recording. Loading a new recording while you have unsaved work asks whether to discard it.

Ellipsis style

The team writes a pause and a broken-off word with a single ellipsis character …. If you prefer three dots `...`, change it in the settings (Notation tab). The validator understands both forms, but shows ellipses written differently from the chosen form as hint VAL-05. The **Unify ellipses** button above the warning list replaces them all at once — a single `Ctrl+Z` undoes the whole replacement. This is particularly useful for texts written in an older version of the application. The note `[...]` and the marker `[…?]` have a fixed form and are not changed.

Find and replace

`Ctrl+F` opens the search bar below the editor; matches are highlighted (on screen only — the highlight does not go into the DOCX), and a counter shows which one is current. Options: **Aa** — match case, **Whole words**, **.*** — regular expression. An invalid expression does not damage the text: the bar shows a message.

`Ctrl+H` adds the **Replace with** field. The ^ character makes the next letter superscript: replacing `będzie` with `be^ndzie` gives be^{n}dzie (a preview is shown next to the field). Type a literal ^ as \^. **Replace all** is a single undo step. The replacement is always literal, even with regular expressions.

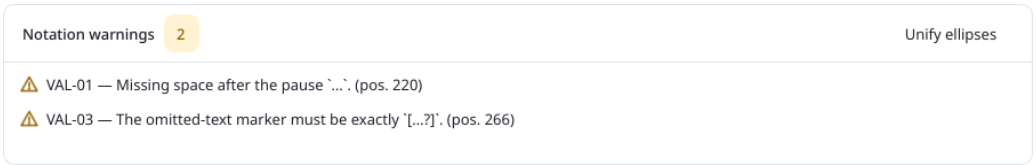
A typical workflow

1. In the header of the **Transcription** card type the **Author**; check the **Date** (today by default, with a drop-down calendar).
2. Click **Import recording…** and choose the recording. The reference code in the recording details is suggested from the file name; fill in the remaining fields (chapter 7).
3. Save the project right away (`Ctrl+S`) — preferably next to the recording.
4. Play short fragments (`Ctrl+Space` or `F4`) and type them in the editor; `Alt+←` skips back a few seconds. If needed, slow down the speed or loop a difficult fragment.
5. Mark pronunciation phenomena: select nasal letters and an inserted *u* and press `Ctrl+Shift+↑` ; insert pauses and markers with the shortcuts.
6. Keep an eye on the **Notation warnings** card below the editor.
7. Save the project from time to time (`Ctrl+S`). The next day open it (`Ctrl+O`) — you will be back at the same position in the recording.
8. Finally: **Export DOCX…** (`Ctrl+E`).

Notation warnings

The application checks the notation as you type. Places that may be inconsistent with the notation are underlined in the text with a wavy line — orange (warning) or grey (hint); hover over it to see the description. The same warnings are listed on the card below the editor; the number next to the title says how many there are (green — no warnings). Click a warning to select the fragment in the text. The underlines do not go into the DOCX. Warnings never fix anything by themselves and do not block the export. Drag the border between the editor and the warnings to enlarge the list.

If the notation at a given place is intentional, right-click the underline (or the warning in the list) and choose **Ignore here**. Only that one warning disappears — other occurrences of the rule are still reported — and the decision is saved in the project file. To turn a whole rule off, use the Settings (**Validator** tab, chapter 9).



The Notation warnings card

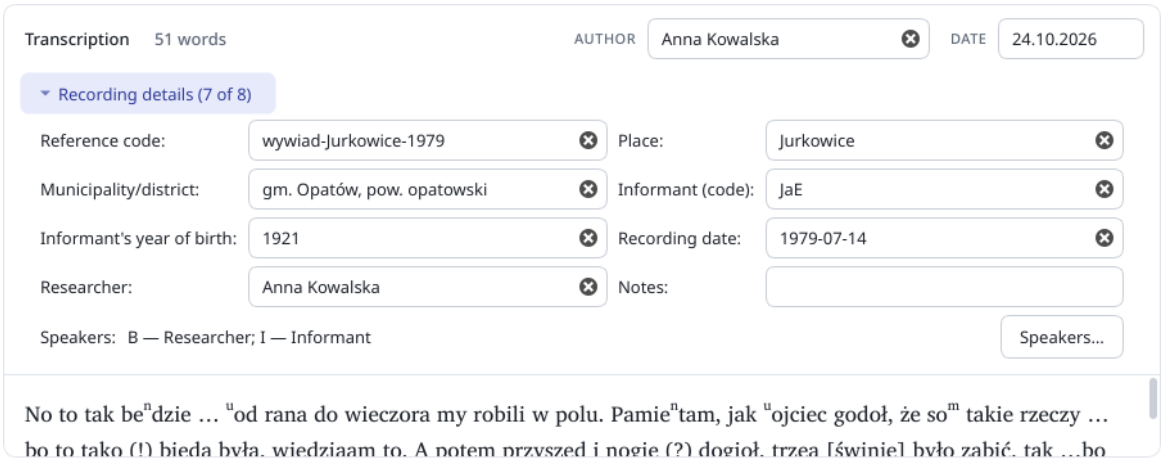
Code	Meaning	How to fix
VAL-01	Pause without a space after the ellipsis: `tak …bo`	add a space: `tak … bo`
VAL-02	Note without a closing bracket: `[o żniwach`	close it: `[o żniwach]`
VAL-03	Malformed omission marker, e.g. `[…]`	use exactly `[…?]` (`Ctrl+3`)

Code	Meaning	How to fix
VAL-05	Hint: an ellipsis written differently from the form chosen in the settings	**Unify ellipses** or fix it manually
VAL-06	Marker stuck to the next word: `nogie (?)dogioł`	add a space after the marker
VAL-07	Hint: several spaces in a row	keep a single space
VAL-08	Hint: space before the closing bracket: `[świnię ]`	remove the space: `[świnię]`
VAL-09	Note inside a note: `[a [b] c]`	split the notes or remove the inner bracket
VAL-10	Superscript on a letter other than *m*, *n*, *u*	select the letter and press `Ctrl+Shift+↑` to remove it
VAL-11	Hint (off by default): *m*, *n* or *u* that may mark nasality or labialization, without superscript	check by listening; if the notation is right — ignore it
VAL-12	Hint: a paragraph without a speaker in a document where other paragraphs have speakers	assign a speaker (`Alt+Shift+1` ...) or ignore it

An ellipsis attached to a word (`jeż…` `kam…` or `jeż...` `kam...`) marks a broken-off word and is not reported.

7. Recording details

The recording details are descriptive data about the recording: reference code, place, informant, recording date and researcher. They are saved in the project and go into a table at the beginning of the DOCX file. Expand them by clicking **Recording details** below the header of the Transcription card; next to it you can see how many fields are filled in. The transcription author and date fields stay on top.



Expanded recording details

Field (default)	Notes
Reference code	Suggested from the recording file name on import when the field is empty (`AdK_1954.aac` → AdK_1954). In the DOCX it becomes the document title.
Place, Municipality/district	The place also goes into the DOCX keywords.

Field (default)	Notes
Informant (code), Informant's year of birth	**Personal data** — left out of the anonymized export.
Recording date, Researcher, Notes	Free text.

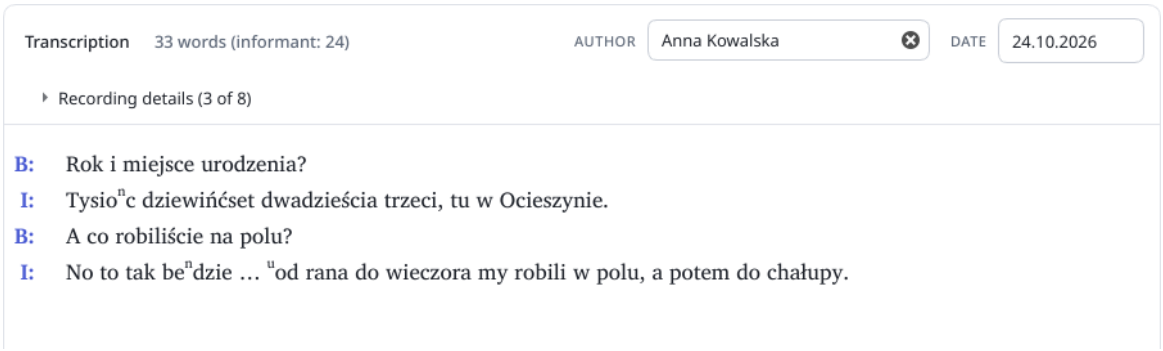
Your team's own fields

In **Settings** (`Ctrl+,`) the **Recording details** tab shows the list of fields. Uncheck a field to hide it in the form and in the export; **Move up** / **Move down** change the order; the **Personal data** column decides about anonymization. **Add field…** creates your own field (e.g. tape number) — rename it with a double-click, and **Remove field** removes it from the list. Default fields cannot be removed, only hidden. Values of hidden fields are not lost — they stay in the project.

The names of the default fields follow the interface language; the names of your own fields stay as you typed them.
During ASR transcription the recording details, like the whole window, are locked.

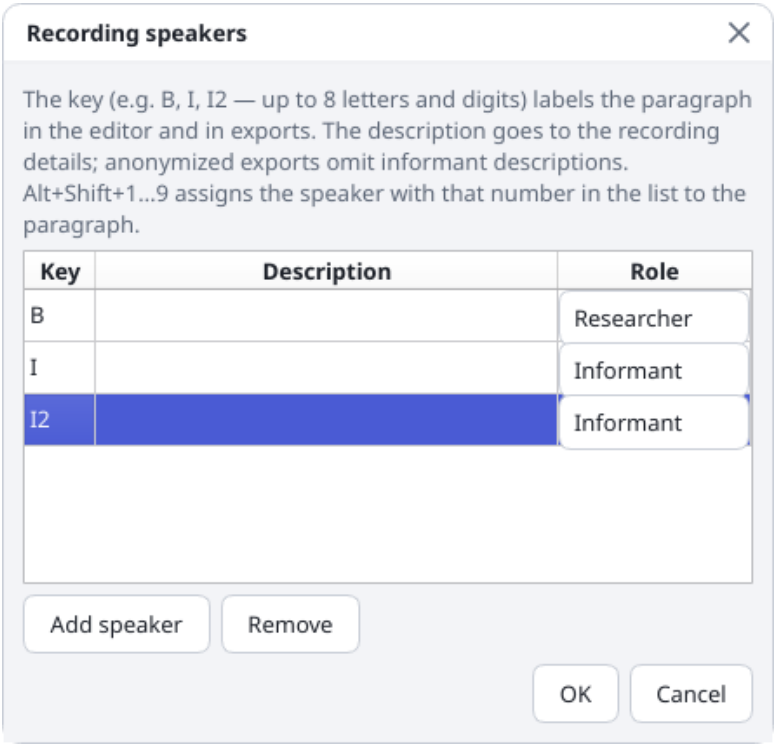
Speakers: researcher and informant

An interview consists of utterances of speakers. Every line of text (paragraph) can have a speaker: put the cursor in the paragraph (or select several) and press `Alt+Shift+1` — the first speaker in the list, by default **B** (researcher), or `Alt+Shift+2` — **I** (informant). Pressing it again removes the assignment, and `Alt+Shift+0` removes the speaker. The same is available in the editor context menu (**Speaker**) and by clicking the label. The label `B:` is shown to the left of the paragraph but is not part of the text — you cannot delete it by accident, and `Ctrl+Z` undoes the assignment in one step. A new paragraph (Enter) gets the speaker of the paragraph it was created in.



Paragraphs of the researcher (B) and the informant (I)

The **Speakers…** button in the expanded recording details edits the recording's speaker list: key (e.g. I2 for a second informant, up to 8 letters and digits), description (e.g. the informant code) and role — researcher, informant or other. Changing a key moves the assigned paragraphs, and removing a speaker who has paragraphs needs confirmation. Speakers are saved in the project, and the legend (*B — Researcher; I — Informant*) goes into the recording details of the export.



The recording's speaker list

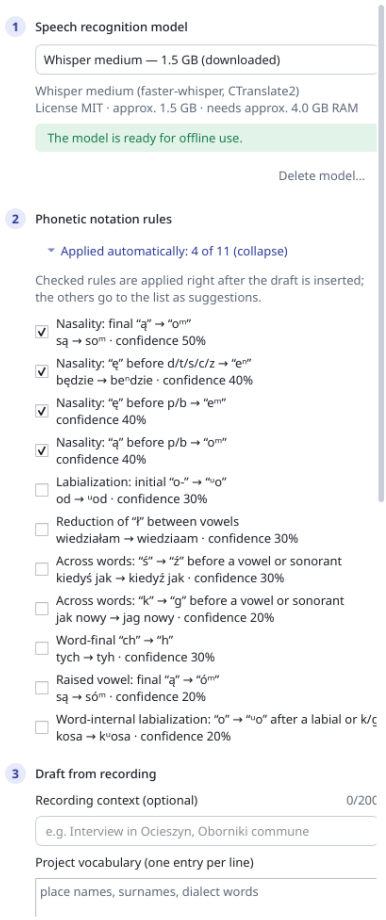
When the document has speakers, the word count also shows the informants' words, the search bar has an **Informant only** box, and the validator points out paragraphs without a speaker (VAL-12). In an anonymized export informant descriptions are replaced by the role name.

8. ASR draft (automatic speech recognition)

The ASR draft is an **optional** helper: the Whisper speech recognition model listens to the recording and produces ordinary (orthographic) Polish text, which goes straight into the editor. Selected phonetic notation rules are applied automatically, and every change is highlighted for your review. Without the model the application works normally.

> **Remember.** The draft is an orthographic, not a phonetic, record. The model does not hear nasality, labialization or the reduction of “ł”, and dialect speech is recognized worse than standard Polish. The rules only guess — every fragment needs your correction.

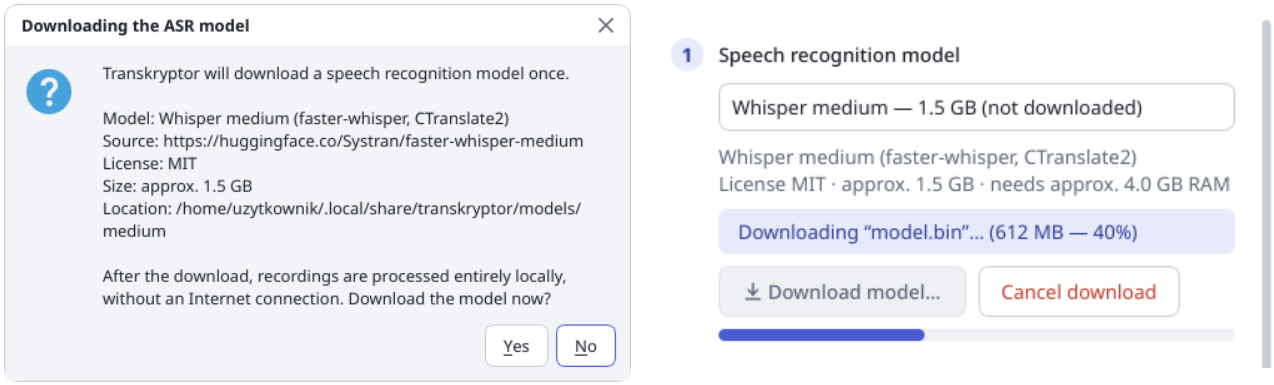
The panel guides you through three numbered steps: model, rules and draft. Below them are the list of segments and the list of changes to review.



Steps 1–3 in the ASR draft panel (rule list expanded)

Step 1 — choosing and downloading the model (once)

1. In the **ASR draft** panel choose a model from the list: **small** (approx. 0.5 GB, fast, for weaker computers), **medium** (approx. 1.5 GB, default) or **large-v3** (approx. 3.1 GB, most accurate, but slower and needs approx. 6 GB of RAM). The choice is remembered. Then click **Download model…**.
2. Read the information: model name, source, license (MIT), size and location. Click **Yes** to download or **No** to cancel.
3. Watch the progress bar. **Cancel download** stops the operation; after an error you can try again — an incomplete download is never treated as a ready model.
4. Once you see the green message **The model is ready for offline use.**, the Internet is no longer needed.



Download consent (left) and download in progress (right)

Each model is downloaded separately. Remove a model you no longer need with **Delete model…** — the application asks for confirmation and shows how much space was freed; other models keep working. If the selected model needs more memory than the computer has, a warning appears before transcription.

On a computer with an NVIDIA graphics card the draft can be created many times faster. This needs the CUDA 12 and cuDNN 9 libraries, installed by you (they are not bundled with the application). Without them the draft is computed on the processor, and a note appears below the transcription button. In **Settings** (**ASR** tab) you can force the processor only.

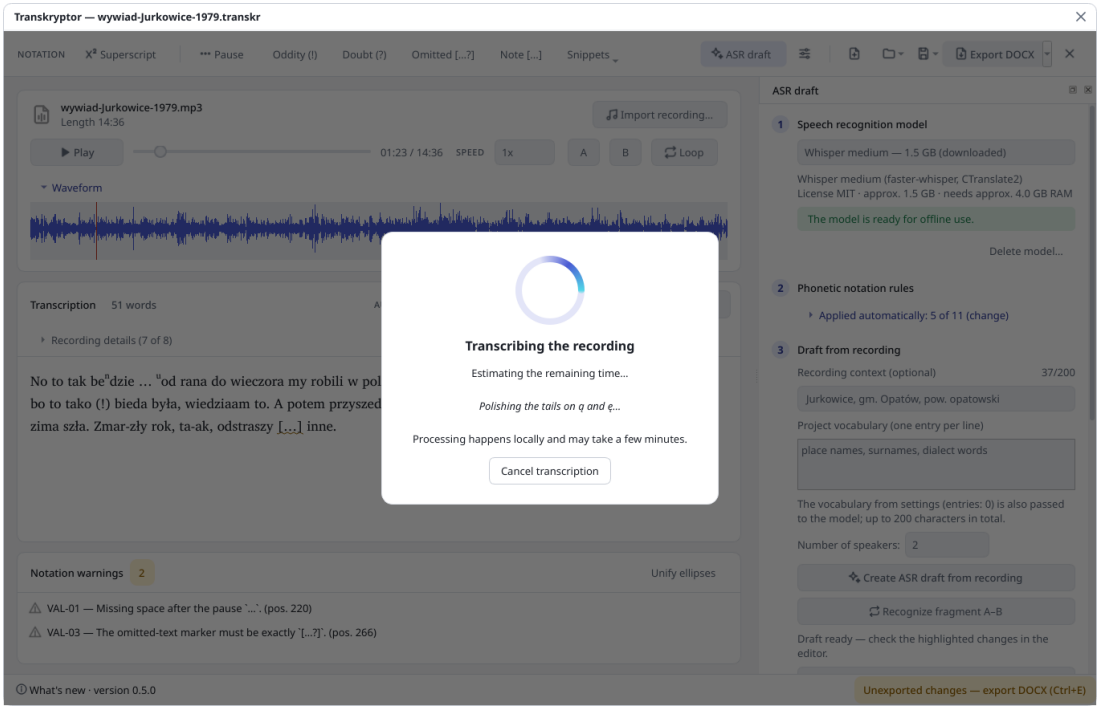
Step 2 — choosing the phonetic notation rules

The rule list is collapsed by default — you only see how many rules are turned on. Click **Applied automatically: … (change)** to expand it. Each rule shows an approximate confidence. **Checked** rules are applied to the draft automatically; matches of **unchecked** rules appear in the list as suggestions for you to decide on. You can change the selection before every draft. The list also contains rules added in the Settings (chapter 9); new built-in rules (e.g. "ɬ" reduction, voicing across word boundaries) are unchecked by default — they only give suggestions.

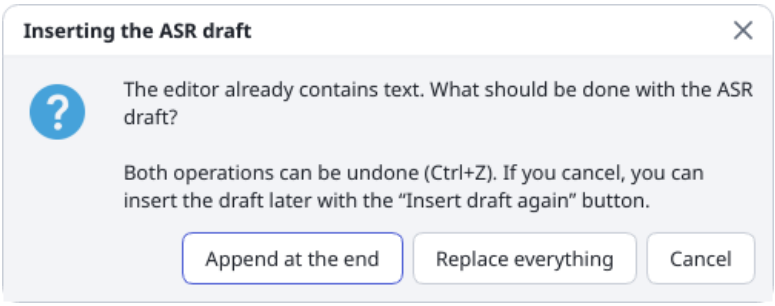
Step 3 — creating the draft

1. Import the recording (**Import recording…** button in the player).
2. Optionally fill in **Recording context** (up to 200 characters; filled in automatically from the place and commune in the recording details) and **Project vocabulary** — place names, surnames, dialect words, one per line. They help the model spell proper names; the project vocabulary is saved in the project file.
3. Click **Create ASR draft from recording**. The button is active when the model is downloaded and a recording is loaded.
4. During the transcription the window is locked and a progress card appears in the middle. Stop it with **Cancel transcription** — the current fragment will be finished first.
5. When it is done, the draft goes into the editor. If the editor already contained text, the application asks whether to **Replace everything** or **Append at the end**; **Cancel** leaves the text unchanged.

> **Careful with hints.** A short, factual context helps, but a long or misleading one may make the model repeat it in the draft. If text from the context that is not in the recording shows up in the draft, shorten or clear the field.



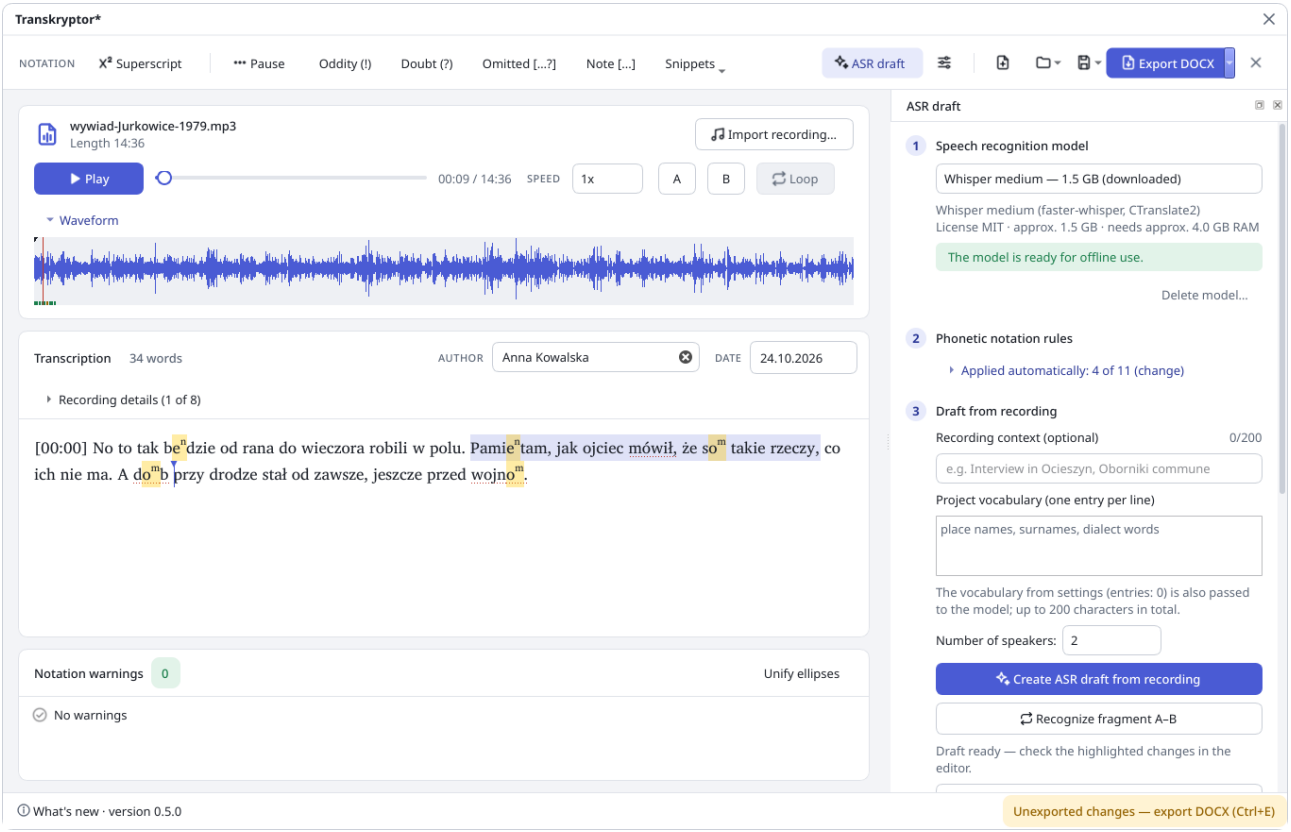
The window during transcription — only cancelling is active



Question about where to put the draft when the editor contains text

Reviewing the changes

Changes made by the rules are highlighted in yellow in the editor. The **Changes and suggestions to review** list shows every match: the word before and after the change, the kind of phenomenon and the confidence.



Draft inserted into the editor — highlighted changes and the review list

Element	Action
✓ (green)	Change applied automatically — highlighted in the editor.
○	Suggestion of an unchecked rule — the text has not been changed.
Clicking an item	Selects the fragment in the editor so that you can easily listen to it and correct it.
Double-click / Apply suggestion	Applies the ○ suggestion to its fragment. If you have changed the fragment since the transcription, the application changes nothing and asks you to correct it manually.
Finish review	Removes the highlights and the list; the text stays unchanged.
Insert draft again	Inserts the last draft once more without repeating the transcription — also after reopening the project.
Segments and confidence	Draft fragments with their time (mm:ss–mm:ss) and approximate confidence. Low-confidence segments are highlighted in colour — listen to them especially carefully. Clicking a segment (or Enter) plays the recording from its start. After importing a different recording the list is cleared.

Undo. Inserting the draft and applying the rules are two separate steps: the first `Ctrl+Z` restores the raw draft without the rules, the second — the state before the insertion.

An unfinished review, the segment list and the last draft are saved in the project: when you open it, the highlights and the list come back and you can keep applying suggestions.

Uncertain words

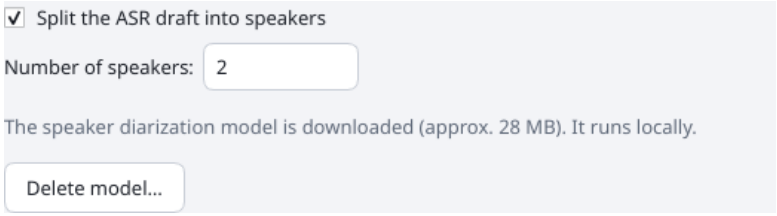
Words the model is unsure about are underlined with red dots in the draft (notation warnings have a wavy line). Hover over a word to see its confidence. `Ctrl` + click on a word plays exactly that word in a loop — pausing or seeking ends the loop. Once you correct the word, the underline disappears. Change the threshold (40% by default) in the settings (**ASR** tab); Whisper confidences are approximate, so the underline is a hint where to listen carefully, not a verdict.

Recognizing a fragment

A difficult passage can be recognized on its own, without repeating the whole recording: mark it with the A–B loop (the A and B buttons or by dragging on the waveform), put the cursor in the text and click **Recognize fragment A–B**. The result goes to the cursor position (or replaces the selected text) with the rules and the link to the recording time; a single `Ctrl+Z` removes it completely. The draft of the whole recording and the segment list stay unchanged. The window is locked during recognition, as for the whole recording.

Splitting the draft into speakers

The draft can be split into speaker turns right away. The first time, click **Download speaker model…** in the **Settings** (**ASR** tab) and confirm the download (approx. 28 MB, Apache-2.0 license) — after that the model runs entirely locally. The **Split the ASR draft into speakers** option is on by default; without the model it is locked and the draft is split into lines every 30 seconds.



Speaker diarization on the ASR tab

After transcription the program recognizes the voices: every turn becomes a separate paragraph with a `[mm:ss]` time marker and a speaker (chapter 7). The first voice in the recording gets the first speaker on the list (**B** by default), the second one **I**, further ones I2, I3… The split also catches quick exchanges of questions and one-word answers (“Occupation?” — “Farmer.”) and places the speaker change at the end of a sentence.

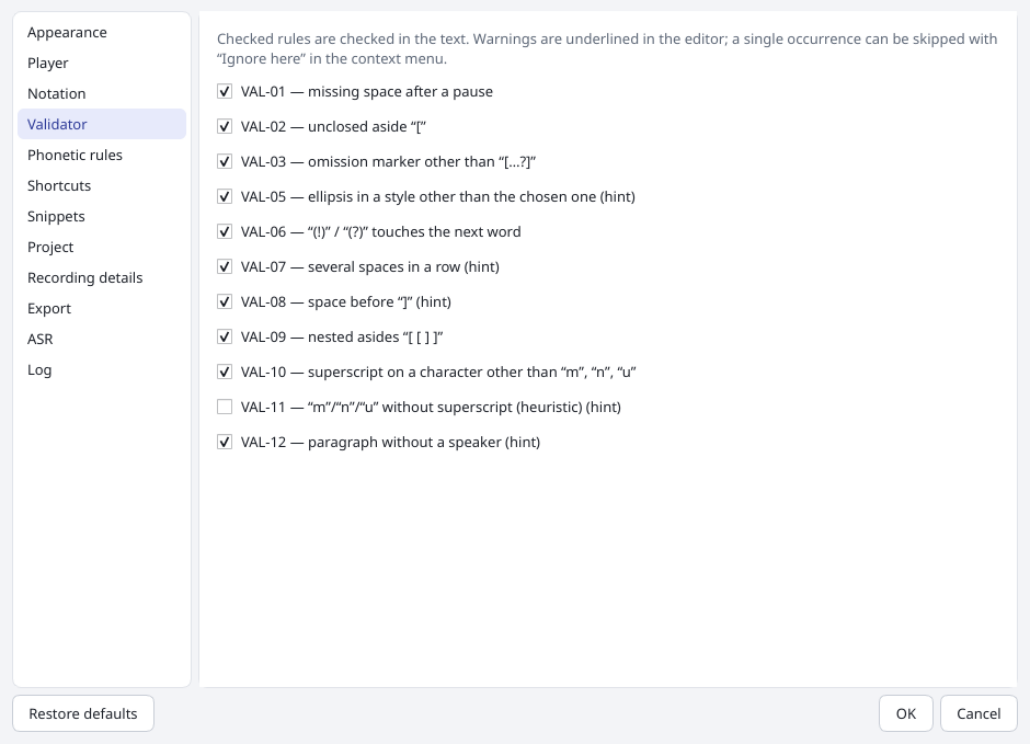
Number of speakers tells the program how many speakers can be heard in the recording (2 by default — researcher and informant, 2–8 to choose from). A given number splits more reliably than **Automatic**, where the program estimates the number of voices itself and may find too many or too few. When splitting is on, the same field is in the **ASR draft** panel (step 3, above the transcribe button), so you can quickly change the number of speakers before each recording — the choice is remembered. Splitting adds a few minutes per hour of recording to the transcription. Recognizing an A–B fragment does not split the text into speakers.

Check the speakers. The split is automatic. It most often errs on very short interjections (e.g. answers while the recording details are read out) and on overlapping voices. Correct the assignment as usual: `Alt+Shift+1` … `9` or the **Speaker** menu.

9. Tailoring the notation: validator, rules, shortcuts

Transkryptor can be adapted to your team's conventions and to the dialect you study. Everything below is set in **Settings…** (`Ctrl+,`); changes apply as soon as you click **OK**.

Validator

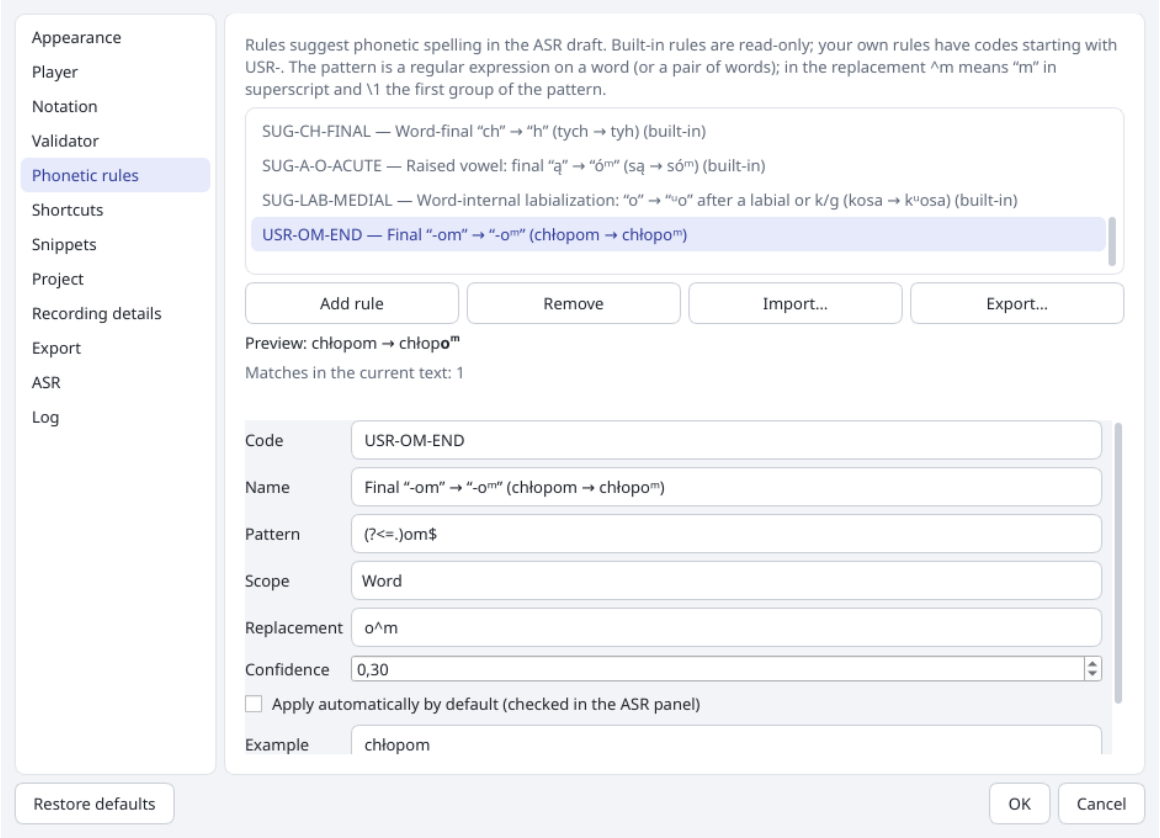


The Validator tab

Checked rules are checked in the text (described in chapter 6). Hints rank lower than warnings. Hint VAL-11 (letters *m*, *n*, *u* without superscript) is off by default because it often reports correct notation — turn it on for the final proofreading if it helps.

Phonetic rules

Rules suggest phonetic spelling in the ASR draft (chapter 8). Besides the built-in rules you can add your own — for example for phenomena typical of the dialect you study. Built-in rules are read-only.



The Phonetic rules tab

- Click **Add rule**. A new rule gets the code `USR-1`, `USR-2`… (the code must start with `USR-`).
- **Pattern** is a regular expression matched against one word, or — with the **Pair of adjacent words** scope — against two words separated by a space (cross-word phonetics).
- **Replacement** is inserted in place of the match: `^m` means *m* in superscript, `\1` the text of the first bracketed group of the pattern.
- **Confidence** is your estimate of how often the rule is right; **Apply automatically by default** checks the rule in the ASR draft panel.
- Type an **Example** — above the form you will see what the rule does and how many matches it has in the current text.

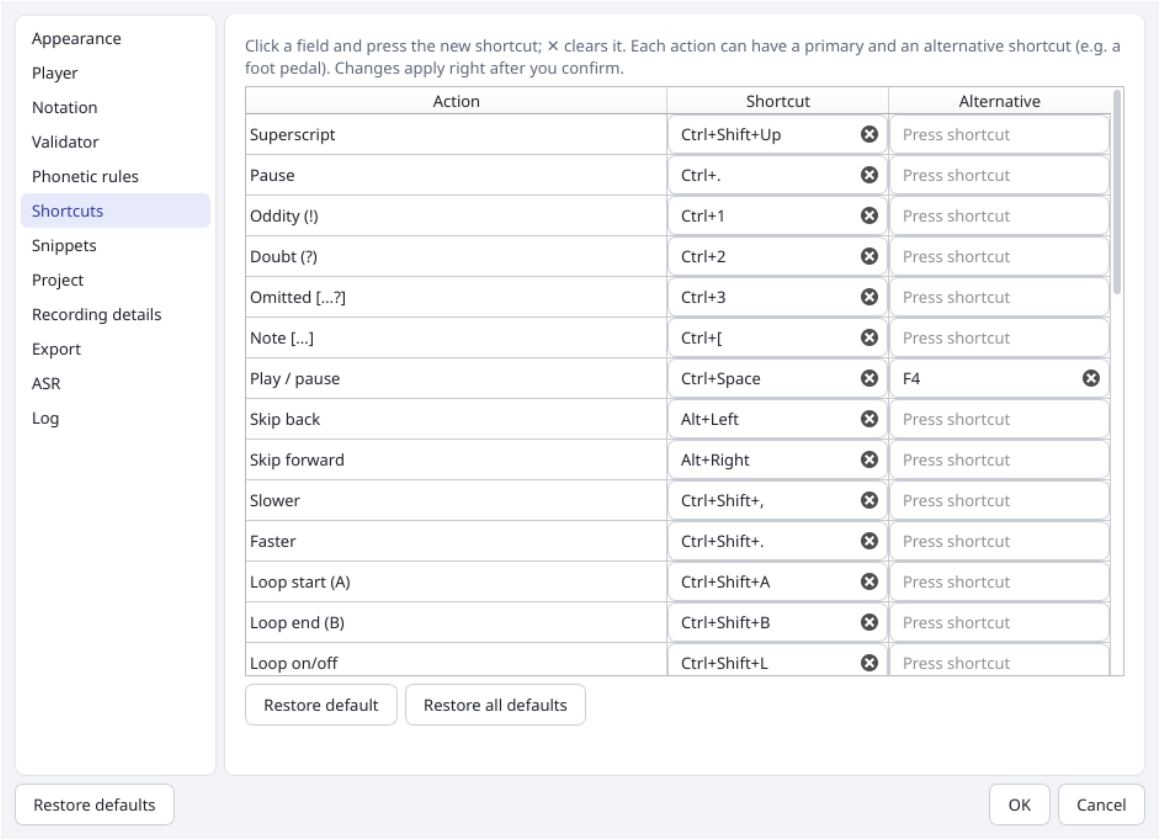
Pattern	Replacement	Effect
`(?<=.)om$`	`o^m`	final *-om*: chłopom → chłopo^{m}
`ę(?=[dtscz])`	`e^n`	*ę* before d/t/s/c/z: będzie → be^{n}dzie
`^o(?=.)`	`^uo`	initial *o-*: od → ^{u}od
`([aeiouy])ł(?=[aeiouy])`	`\1`	*ł* between vowels disappears: wiedziałam → wiedziaam
`ś(?=\s+[aeiouy])` (word pair)	*ź*	kiedyś jak → kiedyź jak

Regular expression cheat sheet: `^` — start of the word, `$` — end of the word, `[pb]` — one of the letters, `(?=…)` — "followed by" (not replaced), `(?<=…)` — "preceded by", `\s+` — a space, `.` — any letter. Letter case matters.

Save your own rule set with **Export…** to a `.toml` file and pass it to other researchers — they add it with **Import…**. If the rules file contains an error, the application shows at start-up which rule and which field are wrong; the other rules keep working. A rule whose pattern takes too long to match is stopped and skipped with a message — the window does not freeze.

New built-in rules Reduction of “ł” between vowels, voicing across word boundaries (kiedyś jak → kiedyź jak, jak nowy → jag nowy), word-final “ch” → “h”, final “ą” → “ó^{m}” and word-internal labialization have approximate confidence and give only suggestions by default. When two rules apply to the same fragment, the rule checked for automatic use wins.

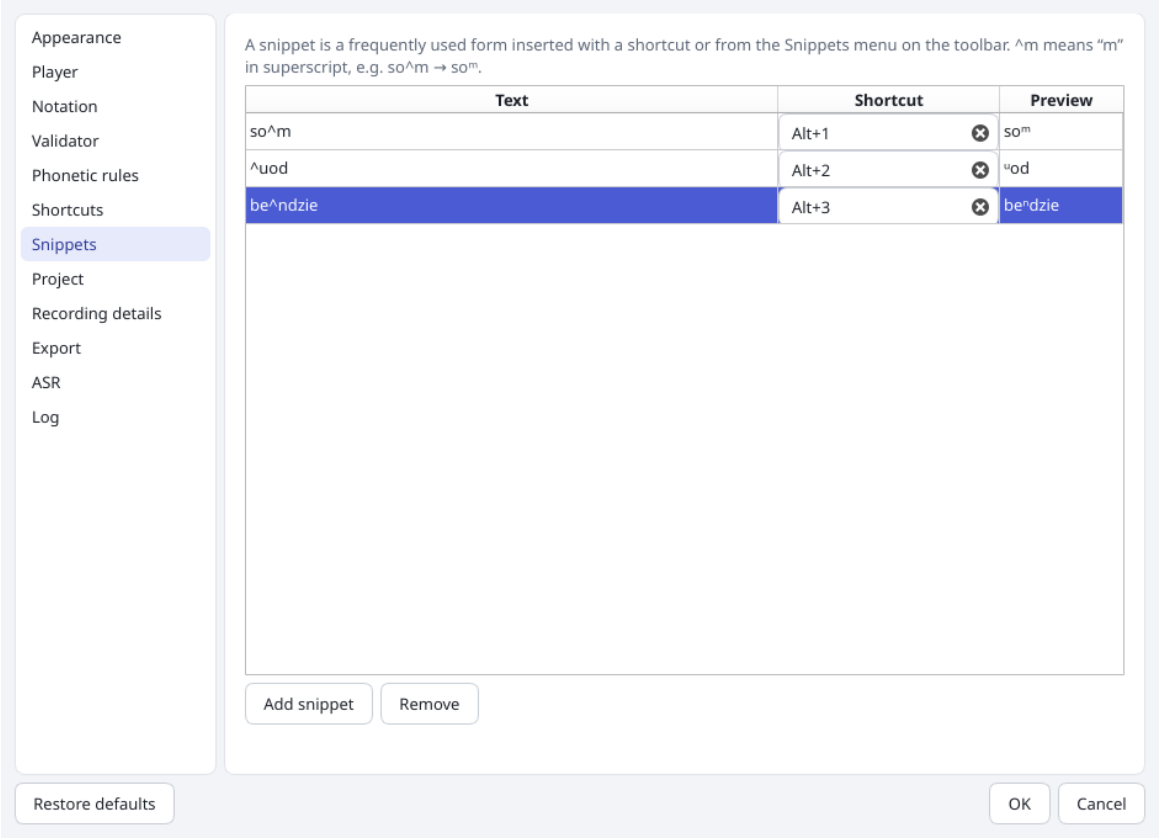
Keyboard shortcuts



The Shortcuts tab

Click a field in the Shortcut column and press the new shortcut; × clears it. Each action can also have an alternative shortcut — useful e.g. for a foot pedal sending F9–F12. If the shortcut is already used by another action or by text editing (copy, paste, undo, word jump…), a description of the conflict appears below the table and the settings cannot be saved until you resolve it. **Restore all defaults** brings back the shortcuts from this guide. Button tooltips always show the current shortcuts.

Snippets and autocomplete



The Snippets tab

A snippet is a frequently used form, e.g. `so^m` (so^{m}) or ^uod (^{u}od). New snippets get a suggested shortcut `Alt+1` … `Alt+9`; you can change or remove it. Snippets are also in the **Snippets** menu on the toolbar. A single `Ctrl+Z` removes the whole snippet.

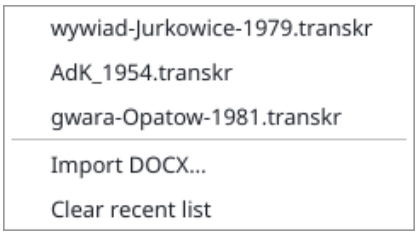
With **Autocomplete** turned on (**Notation** tab) the editor suggests matching snippets and words already used in the text after two letters. `Enter` or `Tab` inserts the suggestion, `Esc` closes it.

10. Project and autosave

The project file (**.transkr**) keeps all the work on a recording: the text with superscripts, the author, date and recording details, a reference to the recording, the playback position and A–B loop, the last ASR draft with its segments and an unfinished rule review. The recording itself is not copied into the project — the project remembers where it is.

Saving and opening

Action	How
Save project	`Ctrl+S` or the floppy disk button. The first save asks for a name and folder (the suggested name is the reference code).
Save project as…	`Ctrl+Shift+S` or the arrow next to the floppy disk — a copy in a new file; you continue working in the new file.
Open project…	`Ctrl+O` or the folder button.
Recent projects	The arrow next to the folder — up to 8 recently saved or opened projects.



The Open project button menu

Keep the project next to the recording. The project remembers where the recording is relative to itself. When you move the whole folder (also to another computer, or from Windows to Linux), the project still finds the recording.

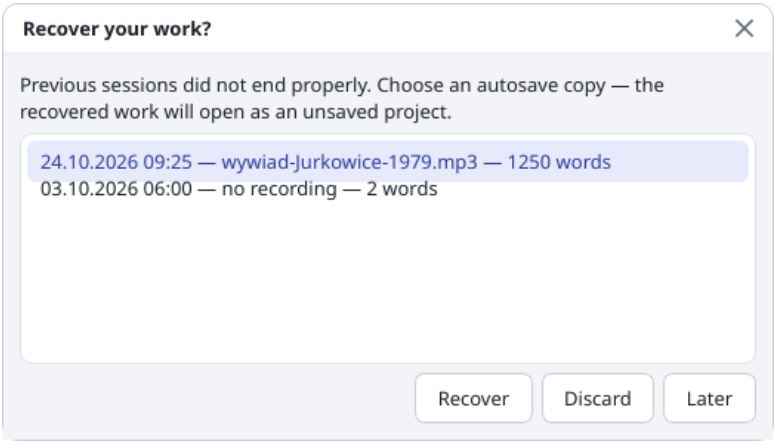
When the recording is missing or different

- **Recording moved or deleted** — the application asks whether to locate the file. **Yes** opens a file dialog, **No** opens the project without the recording: you can keep editing the text and import the recording later.
- **Recording changed or replaced** — the application compares the file checksum. If it does not match, it warns that the segment times and the position may not fit, and asks whether to use the recording anyway.

Autosave

Every 60 seconds (change the interval in the settings, Project tab) the application saves a copy of your current work in its data folder — only when something has changed. Autosave does not change your project file and does not remove the asterisk from the title. During ASR transcription it waits until the end. After a clean exit the copy is deleted.

If the application ends abnormally (a freeze, a power cut), the next time you start it you will see a window with a list of copies: date, recording and number of words.



The recovery window after a crash

Button	Effect
Recover	Opens the selected copy as a new, unsaved project. Save it right away (`Ctrl+S`) — e.g. in place of the previous project.
Discard	Deletes the selected copy permanently.
Later	Keeps the copies — the window appears at the next start.

11. Export: DOCX and other formats

1. Click **Export DOCX…** or press `Ctrl+E` .
2. Choose the folder and the file name (suggested: the reference code from the recording details or `transcription-YYYY-MM-DD.docx`). The .docx extension is added automatically.
3. Click **Save**. The status bar shows “No unexported changes”.

What the file contains

- a recording details table — only the filled-in fields, in the order from the settings, with the speaker legend at the end,
- an *Author: … Date: …* paragraph (if they are filled in),
- the transcription text with all Polish characters, hyphens and markers — every line as a separate paragraph; a paragraph with a speaker starts with a bold label **B:**,
- superscripts as Word formatting (superscript), not as separate characters,
- in the document properties: the author, the title (the reference code or “Phonetic transcription”) and keywords (the place).

The table labels, the *Author/Date* paragraph and the default title are written in the interface language. Transkryptor imports documents exported in either language.

The ASR review highlights do not go into the file. Open the file in Microsoft Word, LibreOffice Writer or Google Docs. If saving fails (e.g. no permission for the folder), you will see a message with a hint — choose a different location and try again.

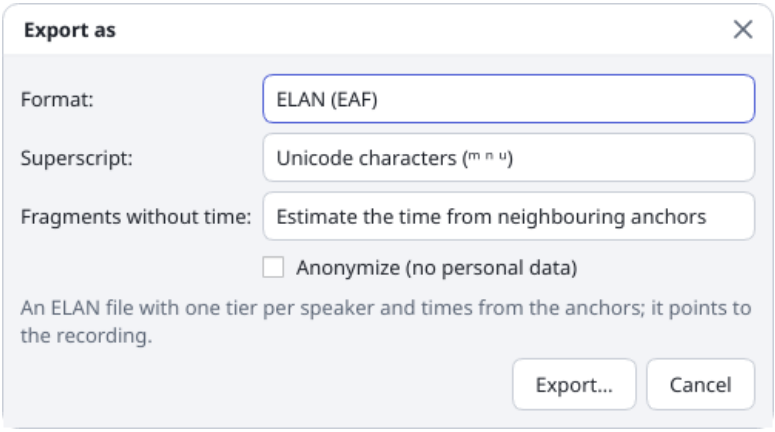
Anonymized export

The arrow next to the **Export DOCX…** button → **Export anonymized DOCX…** creates a file without the recording details fields marked as personal data (by default the informant's code and year of birth) — also in the document properties. The text and superscripts are the same as in a normal export. The suggested file name ends with `-anonymized.docx`. Choose which fields are personal in the Recording details tab of the settings.

> **Note.** Anonymization only applies to the recording details and speaker descriptions. You have to remove names of people or places from the transcription text itself manually.

Other formats: Export as…

The arrow next to the **Export DOCX…** button → **Export as…** opens the format window. Options are shown only for the formats they apply to; your choice is remembered.



The Export as window

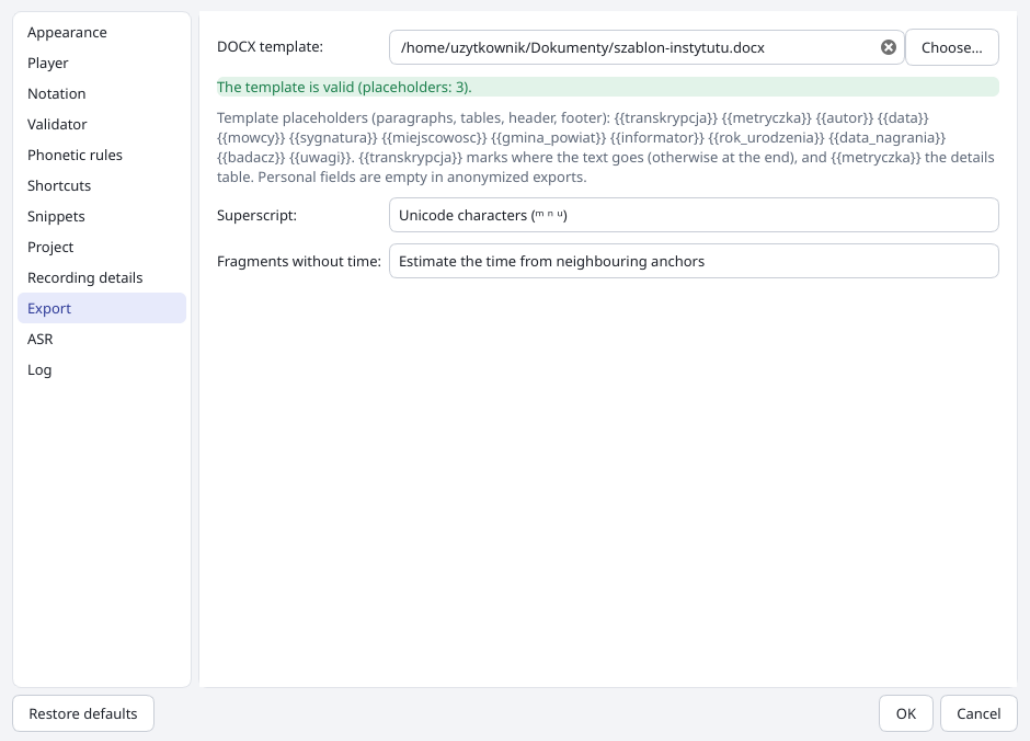
Format	Contents
ODT, PDF	the same layout as DOCX (details, speakers, superscript); PDF A4 with the font embedded
TXT	UTF-8 plain text: details in the header, paragraphs with the prefix `B:`
ELAN (.eaf)	one tier per speaker, times from the anchors, a link to the recording and the details as properties
Praat TextGrid	one interval tier per speaker, times from the anchors
SRT, WebVTT	subtitles with times from the anchors and speakers

Plain-text formats (TXT, ELAN, Praat, subtitles) have no formatting, so superscript is written as **Unicode characters** (`ᵐ ⁿ ᵘ`, the default) or with **caret notation ^m** (a longer fragment: `^{mn}`). A character without a Unicode equivalent is written as ^x and listed in the report.

ELAN, Praat and subtitles need text linked to the recording — an ASR draft or manual timestamps (`Ctrl+Shift+T`); without them they are unavailable. A fragment without time (e.g. typed by hand) gets a time estimated from the neighbouring anchors or is skipped — depending on the **Fragments without time** option. After the export a report lists such fragments.

Institution DOCX template

In **Settings** the **Export** tab lets you choose a .docx file with the required layout (fonts, margins, header, footer, logo). The DOCX export keeps this layout and fills in placeholders in double curly braces: `{{sygnatura}}`, `{{miejscowosc}}`, `{{badacz}}`, `{{informator}}` and so on (the full list is in the tab; English field keys such as `{{signature}}` work too, and case does not matter). `{{transkrypcja}}` marks where the text goes (otherwise it goes at the end), `{{metryczka}}` the details table, and `{{autor}}`, `{{data}}` and `{{mowcy}}` the author, the date and the speaker legend.



The Export tab with a template

- The template is checked as soon as you choose it: the tab shows whether the file opens and whether it has unknown placeholders (those are left unchanged).
- If the template cannot be opened (e.g. it is damaged or is a .doc or .dotx file), the export still succeeds — on the built-in template, with a warning.
- In an anonymized export, placeholders of personal fields are empty.

Font

The default text font is Charis SIL, shipped with the application — it displays phonetic characters and superscript well. PDF embeds the font; DOCX and ODT store its name, so on a computer without Charis SIL the word processor substitutes another font (Charis SIL can be downloaded free of charge from the SIL website). Choose another font in the Appearance tab of the settings; a DOCX template has its own styles.

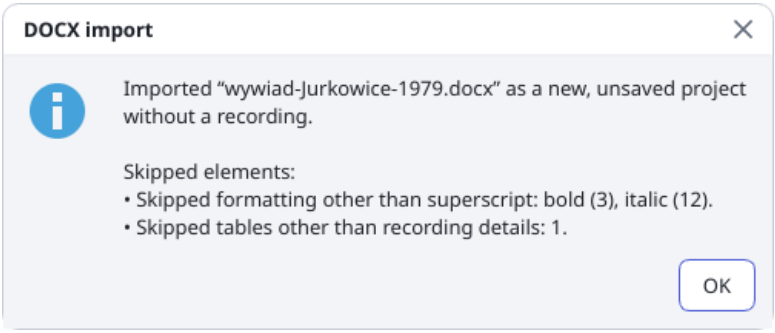
12. DOCX import

You can return to a transcription exported to DOCX earlier (or typed in Word) by importing it: the arrow next to the **Open project** button → **Import DOCX…**. The application creates a new, unsaved project without a recording — import the recording and save the project. If you have unsaved work, it asks first, as with a new document.

What is loaded

- the text and superscripts — every document paragraph as a separate line,
- speakers from bold **B:** labels at the beginning of paragraphs and from the speaker legend in the recording details,
- the recording details from the table at the beginning of the document (recognized by field names in English or Polish, including your team's own fields),
- the author and date from the *Author: … Date: …* paragraph (or its Polish version *Autor: … Data: …*); if it is missing — the author from the document properties.

A file exported from Transkryptor comes back complete. Other documents are loaded on a best-effort basis: paragraphs stay paragraphs, and bold, italics, underline, tables other than the recording details, images and footnotes are skipped. After the import the application shows what it skipped.

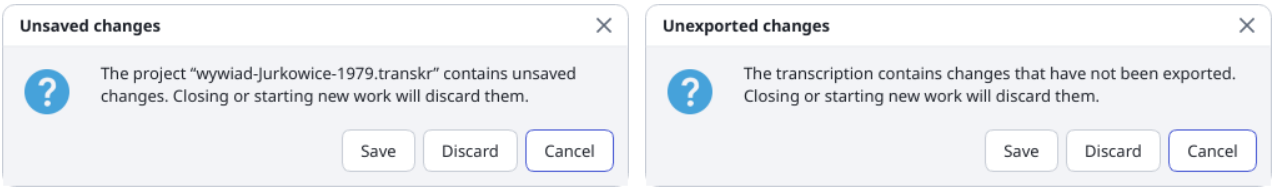


Report after importing a document

13. Closing and unsaved changes

When you close the application, create a new document, load another recording, open a project or import a DOCX, the application checks that no work will be lost:

- **The work has a project file** — it only asks about changes not saved in the project. A saved project that has not been exported to DOCX closes without asking.
- **The work has no project file** — it asks about changes not exported to DOCX.



The question for a project (left) and for work without a project file (right)

Button	Effect
Save	Saves the project (without a file — asks for a name first) and then continues the operation.
Cancel	You go back to work — nothing is lost.
Discard	The changes are discarded permanently and the operation continues.

The application cannot be closed during ASR transcription — cancel the transcription first or wait for it to finish.

14. Privacy

- Recordings and text never leave your computer. There is no cloud speech recognition and no synchronization.
- The application does not collect statistics (telemetry).
- The application connects to the Internet for two purposes only: downloading the ASR model (only after you confirm) and checking whether a newer version is available (at most twice a day, without sending any text, recordings or file paths).
- After the model has been downloaded, everything works offline too.
- You save projects and DOCX files yourself, wherever you choose.
- **Autosave** creates a copy of your work on disk — the text, the recording details (including the informant's code) and the recording path — in the application data folder (table below, `autosave` subfolder). The copy is removed after a clean exit, a new document and saving the project. You can turn autosave off in the settings (Project tab).
- Settings (player, font, language, recording details fields, recent projects list) are stored on your computer: Linux — `~/.config/transkryptor`, Windows — the registry key `HKCU\Software\transkryptor`. They contain no text and no recording details values.
- The **waveform** is remembered in the application cache (Linux — `~/.cache/transkryptor/waveforms`, Windows — `%LOCALAPPDATA%\Transkryptor\cache\waveforms`): only loudness levels every 10 ms, from which speech cannot be reconstructed. You can delete the folder at any time.

Responsibility for the consent of the recorded people and for data protection rules lies with the person or institution doing the transcription.

Where the ASR model and autosave copies are

System	Application data folder
Windows	`%APPDATA%\Transkryptor`
Linux	`~/.local/share/transkryptor`

The ASR model is in the `models` subfolder, autosave copies — in `autosave`.

Remove a single model in the **ASR draft** panel with **Delete model…**. Deleting the whole `models` folder brings back purely manual work and does not affect anything else. The models folder can be moved in the settings (**ASR** tab).

15. Troubleshooting

Problem	What to do
I cannot choose the recording file	The application accepts MP3, WAV, FLAC, M4A, AAC, OGG, OPUS and WMA. Convert other formats, e.g. in Audacity or VLC.
There is no sound	Check the system volume and the output device. On Linux — whether PipeWire/PulseAudio is running.
“Create ASR draft” is greyed out	Download the model and import a recording first.
Downloading the model failed	Check your Internet connection and that you have 3 GB of free space, then click **Download model…** again.
I cannot type during transcription	This is intentional — the window is locked until the transcription ends. Wait or click Cancel transcription.
Transcription takes a very long time	This is normal on a slower processor. You can cancel and try a shorter recording.
A suggestion was not applied	The fragment has changed since the transcription — correct it manually (click the item to select it).
The ASR draft panel disappeared	Click **ASR draft** in the toolbar.
Superscript does not work	Select the letters before pressing `Ctrl+Shift+↑`.
The interface is in the wrong language	**Settings → Appearance → Language**, then restart the application.
Linux: the AppImage does not start	Make it executable (chmod +x), install libfuse2. glibc 2.35+ is required (Ubuntu 22.04 or later).
Windows: SmartScreen warning	More info → Run anyway.
`Ctrl+Space` does not start playback	Use `F4` — on Linux Ctrl+Space is sometimes taken by the input method. The player shortcuts also do not work during ASR transcription.
The project opened without the recording	The recording was moved and “No” was chosen when opening. Click **Import recording…** and save the project.
“Different recording” warning	The recording file differs from the one saved in the project (e.g. it was trimmed or converted). Use it if you are sure it is the same recording.
The application froze or shut down	Start it again and choose Recover (chapter 10).
Bold disappeared after a DOCX import	A transcription only keeps superscript — other formatting is skipped (chapter 12).
Is the application working correctly?	Run it with --self-test (chapter 3).

16. Uninstalling

- **Windows:** Settings → Apps → Installed apps → Transkryptor → Uninstall (this also removes the settings). Then, if you want, delete the `%APPDATA%\Transkryptor` folder (ASR model, autosave copies).
- **Linux:** delete the AppImage file and, optionally, the `~/.local/share/transkryptor` (ASR model, autosave copies) and `~/.config/transkryptor` (settings) folders.

Your projects and DOCX files are not deleted.

Appendix: notation cheat sheet

The phenomena below belong to the Polish dialectological notation used by the application; the examples are Polish words.

Phenomenon	Notation	How in the application
Nasality	be^{n}dzie, pamie^{n}tam, so^{m}	select the letter → `Ctrl+Shift+↑`
Diphthongization, labialization	^{u}od, sw^{u}oje, k^{u}osa	select *u* → `Ctrl+Shift+↑`
Reduction of “ł”	wiedziaam, wyczytaam	write without “ł” — the application does not correct it
Stress group	zmar-zły, ta-ak	a plain hyphen
Modified close vowel	uóna, sóm, słóńce	the letter ó
Pause	rzeczy ... bo	Pause / `Ctrl+.`
Broken-off word	jeż... kam...	ellipsis with no space before it
Pronunciation oddity	zamudze (!)	Oddity / `Ctrl+1`
Uncertain hearing	nogie (?)	Doubt / `Ctrl+2`
Omitted text	odstraszy [...?] inne	Omitted / `Ctrl+3`
Transcriber's note	trzea [świnię]	Note / `Ctrl+[`

Transkryptor is open-source software licensed under GPL-3.0-or-later. The licenses of the libraries and the model are in the `licenses` folder next to the application.